

Initial page

Introducing Pods

Introducing Pods

- · At the heart of Kubernetes
- Most important concept of Kubernetes
- Everything else either manages, exposes or is used by the pods
- · Represents the basic building blocks
- Co-located group of containers
- Common to have one container per pod but not a requirement

Why cant we use the containers directly?

- Containers are designed to run only a single process per container
- If multiple process are running inside a single container it is our responsibility to manage all the process.
- Difficult to manage the logs since all the process with dump the logs on the same container
- · Norm of working with kubernetes is to use single process per container
- Containers in the pod share the same network and hostnames
- Containers in the pod run under the same IPC namespace and communicate via the IPC
- When multiple containers are running, Pod acts as a higher level construct, to manage all the containers together as a single unit. This is the responsibility of a pod
- Pod of containers runs a closely group of containers and run them as if they were running in a isolated environment.
- Best of both the worlds, You get the coherent of related group of process and yet treat individual process as isolated units.
- Kubernetes leverages Docker to run pod
- Pods are other resources can be created by posting the json or YAML file to K8s REST
 API endpoint
- Contains
 - Metadata
 - Specification

Creating a Pod

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4    name: hn-service-pod
5 spec:
6    containers:
7    - image: classpathio/hn-service
8    name: hn-service-container
9    ports:
10    - containerPort: 8011
11    protocol: TCP
```

To get to know the options for creating the pod definition, we can use kubectl explain command

```
kubectl explain pods
kubectl explain pod.spec
```

To create a pod from the pod definition – Generic command to create any resource from yaml/json file

```
kubectl create -f pod-definition.yaml
```

To get the resource description from the resource use the below command

```
kubectl get po hn-service-pod -o yaml
```

```
kubectl get po hn-service-pod -o json
```

Output:

```
1 $ kubectl get pods
2
3 NAME READY STATUS RESTARTS AGE
4 hn-service-pod 1/1 Running 0 6m15s
5 nginx-pod 1/1 Running 0 11m
```

Viewing logs

• Container applications usually log to the standard output and error streams.

Command to check the logs

```
docker logs <container id>
```

Checking the logs with Kubernetes – Used when the pod contains a single container

```
kubectl logs hn-service-pod
```

With multiple containers in a single pod

```
kubectl logs hn-service-pod -c hn-service-container
```

Port forwarding to send requests to the pod from the local machine

```
1 kubectl port-forward hn-service-pod 8111:8111
2
3 Forwarding from 127.0.0.1:8111 -> 8111
4 Forwarding from [::1]:8111 -> 8111
```

Verifying the pod's response

```
1 curl localhost:8111
2 { "hostname":"hn-service-pod" }
```

Organizing pods with labels

- Categorizing pods into subsets
- To operate on every pod belonging to certain group with a single action instead of performing the action on individual pod explicitly
- Done through labels

Introducing labels:

- Organizing resources including pods
- Key-value pairs attaching to resources
- Selecting resources using label-selectors
- Resources can have more than one labels
- Keys should be unique
- Can be created at creation time but can also be modified or add additional labels later without recreating the resource.

Specifying the label during creation time

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
```

```
name: hn-service-pod-label
labels:
env: dev
creation_method: manual
spec:
containers:
- image: classpathio/hn-service
name: hn-service-container
ports:
- containerPort: 8011
protocol: TCP
```

```
1 kubectl create -f hn-service-pod-label.yaml
2
3 kubectl get po --show-labels
```

Listing by label names

```
kubectl get po -L creation_method
  NAME
                        READY
                               STATUS
                                         RESTARTS
                                                   AGE
                                                         CREATION_METHOD
3 hn-service-pod
                        1/1
                               Running
                                                   51m
4 hn-service-pod-label 1/1
                               Running 0
                                                   12m
                                                         manual
5 nginx-pod
                        1/1
                               Running
                                                   56m
```

```
kubectl get po -L creation_method,env
  NAME
                        READY
                                STATUS
                                                   AGE
                                                         CREATION_METHOD
                                         RESTARTS
3 hn-service-pod
                        1/1
                                Running
                                                   51m
  hn-service-pod-label
                        1/1
                                Running 0
                                                   12m
                                                         manual
5 nginx-pod
                        1/1
                                Running
                                                   56m
```

Adding labels after the pod creation

kubectl label po hn-service-pod creation_method=manual notes=test

Overriding existing pod label value

```
kubectl label po hn-service-pod env=prod --overwrite
```

Check if the label value is updated

```
kubectl get po -L env
2 NAME
                           STATUS
                                   RESTARTS AGE
                                                  ENV
                    READY
3 hn-service-pod
                    1/1
                           Running 0
                                             66m
                                                  prod
4 hn-service-pod-label 1/1
                           Running 0
                                             27m
                                                  dev
                    1/1
5 nginx-pod
                           Running 0
                                             71m
```

Label selectors

- We can select the subset of pods and perform action on them.
- They act like filter criteria to shortlist the pods
- Label selectors can be used to select the resources based on
 - Contains or not a label with a key
 - Contains a label with a key and a value
 - Contains the label with a key and not matching value

```
kubectl get po -l creation_method=manual
```

```
kubectl get po -l env
```

```
kubectl get po -l '!env'
```

creation_method!=manual

env in (prod,devel)

env notin (prod,devel)

Use comma seperated label to apply multiple filters

Labels are at the worker node level

To categorize the node. Ex: GPU intensive, SSD's or spinning disks etc

```
kubectl get nodes
                                                 STATUS
                                                          ROLES
                                                                   AGE
                                                                          VΕ
  ip-172-20-43-227.ap-south-1.compute.internal
                                                 Ready
                                                                   115m
                                                                          ν1
                                                          master
                                                                   114m
4 ip-172-20-49-73.ap-south-1.compute.internal
                                                                          ν1
                                                 Ready
                                                          node
5 ip-172-20-70-31.ap-south-1.compute.internal
                                                 Ready
                                                          node
                                                                   114m
                                                                          ν1
6 ip-172-20-73-76.ap-south-1.compute.internal
                                                 Ready
                                                                   114m
                                                                          ν1
                                                          node
```

- 1 kubectl label node ip-172-20-49-73.ap-south-1.comput e.internal gpu=true
- 2 node/ip-172-20-49-73.ap-south-1.compute.internal labeled

```
kubectl get nodes -L gpu
  NAME
                                                 STATUS
                                                          ROLES
                                                                   AGE
                                                                          VΕ
3 ip-172-20-43-227.ap-south-1.compute.internal
                                                 Ready
                                                          master
                                                                   120m
                                                                          ٧1
4 ip-172-20-49-73.ap-south-1.compute.internal
                                                 Ready
                                                          node
                                                                   118m
                                                                          ν1
                                                 Ready
5 ip-172-20-70-31.ap-south-1.compute.internal
                                                                   118m
                                                                          ν1
                                                          node
  ip-172-20-73-76.ap-south-1.compute.internal
                                                                          ٧1
                                                 Ready
                                                          node
                                                                   118m
```

Scheduling pods to a specific node label

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4    name: hn-service-pod-label-gpu
5    labels:
6    env: dev
```

```
7 creation_method: manual
8 spec:
9 nodeSelector:
10 gpu: "true"
11 containers:
12 - image: classpathio/hn-service
13 name: hn-service-container
14 ports:
15 - containerPort: 8011
16 protocol: TCP
```

Annotating pods

- Also key value pair like labels
- Aren't meant to hold information
- Cannot be selected unlike label selector
- Primary meant to be used by tools
- Some annotations are automatically added by the tools
- Used to introduce new features to the specification in a graceful manner
- Once the new features are agreed upon, the old ones are depricated

```
kubectl get po
                                  STATUS
2 NAME
                           READY
                                            RESTARTS
                                                      AGE
3 hn-service-pod
                           1/1
                                   Running 0
                                                      106m
4 hn-service-pod-label
                           1/1
                                   Running 0
                                                      67m
5 hn-service-pod-label-gpu
                                                      9m15s
                           1/1
                                   Running 0
6 nginx-pod
                           1/1
                                           0
                                   Running
                                                      111m
```

Just like label's annotations can be added during creation time or added later

```
1 kubectl annotate po hn-service-pod classpathio/type= "backend API"
2 pod/hn-service-pod annotated
```

Run the describe command to see the annotation

```
kubectl describe po hn-service-pod
                hn-service-pod
   Name:
               default
3 Namespace:
4 Priority:
               ip-172-20-70-31.ap-south-1.compute.internal/172.20.70.31
5 Node:
6 Start Time: Fri, 27 Mar 2020 06:34:34 +0000
7 Labels:
              creation_method=manual
                env=prod
                notes=test
10 Annotations: classpathio/type: backend API
                kubernetes.io/limit-ranger: LimitRanger plugin set: cpu requ
                Running
12 Status:
13 IP:
                100.96.1.3
14 IPs:
    IP: 100.96.1.3
```

Namespaces to group resources

- Grouping of resources can be done using label s
- Labels can overlap since multiple resources can container many labels
- Namespaces are used to split objects into separate non-overlapping groups
- Can be used to split resources based on tenant, environments etc.
- Resource names should be unique withing a namespace
- Works like packages in Java

Operations on a namespace

List all namespace

Fetch the resources specific to a namespace

```
kubectl get po --namespace kube-system
   NAME
                                                                          REA
      RESTARTS AGE
   dns-controller-5769c5f8b6-nhznh
                                                                          1/1
                 172m
   etcd-manager-events-ip-172-20-43-227.ap-south-1.compute.internal
                                                                          1/1
      0
                 172m
8 etcd-manager-main-ip-172-20-43-227.ap-south-1.compute.internal
                                                                          1/1
                 172m
                                                                          1/1
   kops-controller-p4b8n
                 172m
     0
12 [ec2-user@ip-172-31-36-76 pods]$ kubectl get po -n kube-system
13 NAME
                                                                          REA
      RESTARTS AGE
15 dns-controller-5769c5f8b6-nhznh
                                                                          1/1
                 173m
```

i Use -n instead of -namespace

Creating a namespace

```
1 apiVersion: v1
2 kind: Namespace
3
4 metadata:
5 name: classpath-dev-namespace
```

- kubectl create -f classpath-dev-namespace.yaml
- 2 namespace/classpath-dev-namespace created

Alternate way instead of creating from YAML file

- kubectl create namespace classpath-qa-namespace
- 2 namespace/classpath-qa-namespace created

Creating resources inside the Namespacespace

- 1 kubectl create -f hn-service-pod-label.yaml -n "classpath-dev-namespace"
- pod/hn-service-pod-label created
- When performing operationg like listing, deleting we need to specify the namespace
 using the -n or --namespace
- If not specified, the default namespace will be used from the configured context
- The context can be configured to set the default namespace using the kubectl config commands

kubectl config set-context classpath-dev-namespace --namespace

Stopping Pods

Deleting the pod by name

- 1 kubectl delete po nginx-pod
- pod "nginx-pod" deleted
- First the K8s sends the SIGTERM signal to the process and waits for certain amount of time (30 seconds by default)
- If the process does not terminate, SIGKILL signal is issues to kill the process.
- Make sure that the process is shutdown gracefully (shutdown using the command)

Deleting multiple pods

kubectl delete po nginx-pod, hn-service-pod-label

```
pod "hn-service-pod-label" deleted
pod "nginx-pod" deleted
```

Deleting pods using the label selector

```
1 kubectl delete po -l env=dev
2 pod "hn-service-pod" deleted
```

Deleting all the resources including pods by deleting the namespace

```
1 kubectl delete ns classpath-dev-namespace
2 namespace "classpath-dev-namespace" deleted
```

Deleting the pods by retaining the namespace

```
1 kubectl delete po --all
2 pod "hn-service-pod-label-gpu" deleted
```

i) Note: creating a pod using the kubectl run creates a ReplicationController and hence, deleting the pod will not work as the RC will spin up a new pod.

Deleting the pods, rc and the namespace

```
1 kubectl delete all --all
2    pod "hn-service-pod-label-gpu" deleted
3    replicationcontroller "test" deleted
4    service "test" deleted
5    service "test-http" deleted
```

Additional Notes

- First all represents all resources types (pods, rc)
- Second –all represents deleting all resources instead by name
- The command also list all the resources that were deleted
- The service will also get deleted but will be automatically created after sometime.

Replication Controller

Livess Probe

- Pod is scheduled to a Node
- kubelet on the node will run its container
- Keep the containers running as long as the pod exists
- In case of process crash, the kubelet restarts the container
- The process should not restart in case of unhealty pod
- Monitor the health from outside the container and not depend on the app
- Checking container liveness using Liveness Probe
 - HTTP GET probe (2xx and 3xx status codes are considered successfull)
 - TCP Socket (connection successfully established is considered successfull)
 - Exec (Exit status code of 0 is considered successfull)

Creating a Liveness Probe

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4    name: hn-service-unhealthy
5 spec:
6    containers:
7    - image: classpathio/hn-service-unhealthy
8    name: hn-service-unhealthy
9    livenessProbe:
10    httpGet:
11    path: /
12    port: 8111
```

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4    name: hn-service-unhealthy
5 spec:
6    containers:
7    - image: classpathio/hn-service-unhealthy
8    name: hn-service-unhealthy
9 livenessProbe:
```

```
10 httpGet:
11 path: /
12 port: 8111
```

```
kubectl create -f liveness-probe. yaml
pod/hn-unhealthy-service created
```

```
    kubectl get po
    NAME
    READY STATUS RESTARTS AGE
    hn-unhealthy-service
    1/1
    Running
    0
    57s
```

```
    kubectl get po
    NAME READY STATUS RESTARTS AGE
    hn-unhealthy-service 1/1 Running 0 57s
```

Restarts column displays the number of times the pod restarted

Obtaining the logs

```
$\text{$kubectl logs hn-unhealthy-service } --previous$

Inside the getHostname method of Controller

Number of requests 5

Received request from 100.96.1.1You have hit the Server

Your current IP address: hn-unhealthy-service/100.96.1.6

Your current Hostname: hn-unhealthy-service

Inside the getHostname method of Controller

Number of requests 6

Inside the getHostname method of Controller

Number of requests 6

Inside the getHostname method of Controller

Number of requests 6

12 Inside the getHostname method of Controller

Number of requests 6

13 2020-03-27 10:22:02.860 INFO 1 --- [extShutdownHook] o.s.s.concurrent.Thr
```

Configuring additional properties

```
1 livenessProbe:
2 httpGet:
3 path: /
4 port: 8111
5 initialDelaySeconds: 15
```

Points to consider:

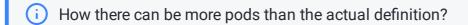
- The Liveness probe should be light
- No business logic in the liveness probe
- · Definition of the healthy service should be defined here
- No retry loops as it is handled by the K8s itself
- Should not take too long to complete
- · Should not be CPU intensive

Replication Controller

- Replication controller is a K8s resource
- Ensures that the pods are always running.
- If a running pod goes missing, the replication controller crates a replacement pod across pods.

Operations of Replication controller

- Continuously monitors the running pods and ensures that the actual number of running pods match the desired number.
- If two many pods are running, the Replication controller, removes the excess pods.



Answer:

Pods created manually

- Changing the pods definition
- Changing the desired number of pods in the template

Operations of a Replication Controller

- ReplicationController work on a set of pods which matches the label
- Replication controller's job is to ensure the actual running pods are matching the labelselector.
- Reconcile to match the desired count
- ReplicationController has three essential parts
 - Label selector matching pods
 - Replica count the desired count
 - Pod template the template of pod (definition)
- All the above entries can be modified at any time, but only the replica count affects the current pods
- Changing the label selectors does not affect the existing pods
- Existing pods fall out of the scope
- Controller stops caring about the pods which are out of scope

Benefits of Replication Controller

- Ensures that the desired number of pods are always running, even in case of pod failure
- Sustains the node failure
- Enables horizontal scaling of pods
 - both manual and automatic

Creating a Replication controller

```
1 apiVersion: v1
2 kind: ReplicationController
3 metadata:
4    name: hn-service
5 spec:
6    replicas: 3
```

```
7  selector:
8    app: hn-service-api
9    template:
10    metadata:
11    labels:
12    app: hn-service-api
13    spec:
14    containers:
15    - name: hn-service-container
16    image: classpathio/hn-service
17    ports:
18    - containerPort: 8111
19    protocol: TCP
```

(i) Caution:

- If the selector does not match the pod selector, the API server validates and will not accept, if it is misconfigured
- K8s will automatically extract the label selector, if it is not part of the replication controller
- Use this approach and do not specify the label selector in the Replication controller

Creating a replication controller

```
$ kubectl create -f hn-service-rc.yaml
```

To verify the rc in action

```
$ kubectl get pods
```

Delete a pod to see if the RC creates a new pod

kubectl delete pod hn-service-5sz5t

```
kubectl get pods
```

Getting RC information

```
kubectl get rc
```

Explain: To get Additional information about the RC

```
kubectl describe rc hn-service
  Name:
                hn-service
  Namespace: default
Selector: app=hn-service-api
6 Labels:
               app=hn-service-api
  Annotations: <none>
                3 current / 3 desired
  Replicas:
  Pods Status: 3 Running / 0 Waiting / 0 Succeeded / 0 Failed
  Pod Template:
   Labels: app=hn-service-api
    Containers:
    hn-service-container:
                  classpathio/hn-service
      Image:
      Port:
                    8111/TCP
     Host Port:
                    0/TCP
     Environment: <none>
      Mounts:
                    <none>
    Volumes:
                    <none>
  Events:
    Type Reason
                                    From
                              Age
                                                            Message
    Normal SuccessfulCreate 7m27s replication-controller Created pod: hn
    Normal SuccessfulCreate 7m27s replication-controller Created pod: hn
    Normal SuccessfulCreate 7m27s replication-controller Created pod: hn
    Normal SuccessfulCreate 4m8s
                                    replication-controller Created pod: hn
    Normal SuccessfulCreate 3m12s replication-controller Created pod: hn
```

i Events should also be displayed at the bottom after running the describe command

Behind the scenes

- RC is notified about the deletion of pod event
- API server allows for clients to watch for changes/events
- The notification triggers the RC to verify the number of pods against the desired count to take appropriate action

Responding to node failure

List the pods with -o wide option to get the address of the node

```
kubectl get node -o wide
```

```
gcloud compute ssh gke-kubia-default-pool-b46381f1-zwko

Enter passphrase for key '/home/luksa/.ssh/google_compute_engine':

Welcome to Kubernetes v1.6.4!

...

sudo ifconfig eth0 down
```

```
kubectl get node
```

Moving the pods' in and out of the Replication Controller

Pods created using the Replication Controllers are not tied up with the RC in any way

- Pod contain a metadata.ownerReferences field which we can use to idendify the replication controller the pod belongs to
- If we change the label of the pod, then it is no more managed by the RC and is as good as creating the pod manually
- The RC will spin a new pod, in that case since the pod definition changed
- Adding another label does not affect the RC as long as the RC matches the pod labels

```
1 kubectl label pod kubia-dmdck type=special
2
3 pod "kubia-dmdck" labeled
```

```
kubectl get pods --show-labels
  NAME
                    READY
                           STATUS
                                    RESTARTS
                                               AGE
                                                    LABELS
  hn-service-4nc6d
                    1/1
                           Running
                                                    app=hn-service-api
                                               34m
4 hn-service-755n6
                                                    app=hn-service-api
                    1/1
                           Running
                                               37m
                                    0
  hn-service-g6zgl
                           Running
                                               33m
                                                    app=hn-service-api
                    1/1
                                    0
```

Example of overriding the existing label value to drift from the rc definition which will notify the RC and RC will take action – In this case, creating a new pod from the template

- kubectl label pod hn-service-4nc6d app=non-existent --overwrite
 pod/hn-service-4nc6d labeled
- kubectl get pods -L app **READY** STATUS RESTARTS AGE **APP** 3 hn-service-4nc6d Running non-existent 1/1 36m 4 hn-service-755n6 1/1 Running 0 39m hn-service-api hn-service-g6zgl 1/1 Running 0 35m hn-service-api hn-service-sl7s9 1/1 Running 23s hn-service-api 0

Can be used to remove a unhealthy pod from the rc scope and deleted later

kubectl edit rc hn-service

This is open the YAML inside a editor which you can open and save the file

After editing the RC template with the new pod definition, delete the pods so that the RC will create all the pods with the new definition – Can be used for updates

Changing the default editor

```
export KUBE_EDITOR="/usr/bin/nano"
```

Scaling Pods

• Scaling the pods using the RC - for experimental purpose only since it is not versioned

```
kubectl scale rc hn-service--replicas=10
```

Editing the file using declarative approach - Can be versioned

```
kubectl edit rc hn-service
```

Verify the RC definition after the changes

```
1 kubectl get rc2 NAME DESIRED CURRENT READY AGE3 hn-service 5 5 51m
```

Scaling down - Modifies the spec.replicas in the RC definition like the edit we did in the RC definition

kubectl scale rc hn-service --replicas=3

Deleting the Replication Controller

- Deleting an RC
 - Delete the RC but keep the existing pods and do not delete them Default is delete the pods also along with the RC
 - kubectl delete rc hn-service --cascade=false
 - 3 replicationcontroller "hn-service" deleted
- i Note: You can manage the pods again by creating a RC to match the pods label again

ReplicaSets

ReplicaSets

- ReplicaSets replaces the RC
- Should be used for all new work
- Use ReplicaSets instead of RC

ReplicaSets vs RC

- Works in the same fashion as RC
- More expressive pod selectors
- RC label selector only allows for matching pods including a certain label (not values of the label key/value)
- RS allows for selector allows matching pods that lack a label or pods which include a certain label key, regardless of the value
- Example:
 - RC cannot match pods with label env=production and env=dev at the same time
 - Match either pods with env=production or env-=dev at a single time
 - RS can match both of them at the same time and treat them as a single group
 - RC can also select the pods with only key regardless of the value. Example env=*

Defining a RS

```
1 apiVersion: apps/v1
2 kind: ReplicaSet
3 metadata:
4   name: hn-service-rs
5 spec:
6   replicas: 3
7   selector:
8   matchLabels:
9    app: hn-service
10   template:
11   metadata:
12   labels:
13    app: hn-service
14   spec:
15   containers:
```

```
16 - name: hn-service-container
17 image: classpathio/hn-service
```

- API is in a different namespace
- Only different is the selector compared to a RC
- Instead of listing labels, the RS has selector.matches property instead of under selector property in RC

Creating a ReplicaSet

```
    kubectl get rs
    NAME DESIRED CURRENT READY AGE
    hn-service-rs 3 3 2m37s
```

Describe Replica Set

```
1 Name: hn-service-rs
2 Namespace: default
3 Selector: app=hn-service
4 Labels: <none>
5 Annotations: <none>
6 Replicas: 3 current / 3 desired
7 Pods Status: 3 Running / 0 Waiting / 0 Succeeded / 0 Failed
8 Pod Template:
9 Labels: app=hn-service
10 Containers:
11 hn-service-container:
12 Image: classpathio/hn-service
13 Port: <none>
```

RS - selector example

```
1 apiVersion: apps/v1
2 kind: ReplicaSet
3 metadata:
4 name: hn-service-rs-match-expression
```

```
spec:
  replicas: 3
 selector:
    matchExpressions:
      - key: app
        operator: In
        values:
          - hn-service-api
 template:
    metadata:
      labels:
        app: hn-service-api
   spec:
      containers:
      - name: hn-service-container
        image: classpathio/hn-service
```



Note:

If multiple expressions are put, all of them should evaluate to true

Listing ReplicaSets

```
    kubectl get rs
    NAME
    besired Current Ready AGE
    hn-service-rs-match-expression
    3
    3
    3
    3
```

Listing Pods

```
kubectl get po
                                               STATUS
                                                         RESTARTS
NAME
                                       READY
                                                                     AGE
hn-service-rs-match-expression-5mfb8
                                       1/1
                                               Running
                                                                     2m33s
                                                         0
hn-service-rs-match-expression-ppnvl
                                       1/1
                                               Running
                                                         0
                                                                     2m33s
                                       1/1
hn-service-rs-match-expression-r6x2b
                                               Running
                                                         0
                                                                     2m33s
```

- 1 kubectl delete rs hn-service-rs-match-expres sion
- 2 replicaset.apps "hn-service-rs-match-expression" deleted

DaemonSet

- Similar to RC and RS
- · Pods are deployed to all the nodes
- Replicas does not matter since it is at the node level
- Selecting to deploy on a subset of nodes can be donesusing node-Selector property of the pod definition.
- Meant for run the system level services

Example

• Defining the DaemonSet yaml definition

```
1 apiVersion: apps/v1beta2
2 kind: DaemonSet
 3 metadata:
4 name: ssd-monitor
5 spec:
6 selector:
   matchLabels:
    app: ssd-monitor
9 template:
10 metadata:
       labels:
        app: ssd-monitor
spec:
    nodeSelector:
     containers:
      - name: main
        image: luksa/ssd-monitor
```

- kubectl create -f hn-daemonset.yaml
- daemonset.apps/healthchecker created

List the DamesonSet

```
1 kubectl get ds
2 NAME DESIRED CURRENT READY UP-TO-DATE AVAILABLE NODE
3 healthchecker 0 0 0 0 0 disk=
```

List all the Nodes

```
kubectl get nodes
```

Add the label to the nodes

```
kubectl label node ip-172-20-84-241.ap-south-1.compute.internal disk=ssd
node "ip-172-20-84-241.ap-south-1.compute.internal" labeled
```

List all the pods

```
    kubectl get po
    NAME READY STATUS RESTARTS AGE
    ssd-monitor-hgxwq 1/1 Running 0 35s
```

Remove the label from the node

- kubectl label node ip-172-20-84-241.ap-south -1.compute.internal disk=hdd
- 2 node/ip-172-20-84-241.ap-south-1.compute.internal labeled

List the Node with -L option

```
kubectl get nodes -L disk
  NAME
                                                  STATUS
                                                           ROLES
                                                                    AGE
                                                                          VER
3 ip-172-20-44-6.ap-south-1.compute.internal
                                                  Ready
                                                           node
                                                                    48m
                                                                          v1.
4 ip-172-20-52-93.ap-south-1.compute.internal
                                                  Ready
                                                           master
                                                                    50m
                                                                          v1.
5 ip-172-20-60-251.ap-south-1.compute.internal
                                                  Ready
                                                           node
                                                                    48m
                                                                          v1.
6 ip-172-20-84-241.ap-south-1.compute.internal
                                                  Ready
                                                                    48m
                                                                          v1.
                                                           node
```

Again list the pods

```
kubectl get poNo resources found in default namespace.
```

Running a single completable task

- In case of single completable task after which the process terminates
- Does not restart once the process completes successfully
- In case of job failure, we can configure to either restart or exit

Job Resource

Creating a JobResource

```
apiVersion: batch/v1
kind: Job
metadata:
name: batch-job
spec:
template:
metadata:
labels:
```

```
9 app: batch-job
10 spec:
11 restartPolicy: OnFailure
12 containers:
13 - name: job
14 image: classpathio/batch-job
```

Create a Job resource from the template

```
1 kubectl create -f hn-batch-job.yaml
2 job.batch/batch-job created
```

Test the JobResource

```
    1 kubectl get job
    2 NAME COMPLETIONS DURATION AGE
    3 batch-job 1/1 9s 10s
```

```
1 kubectl get po
2 NAME READY STATUS RESTARTS AGE
3 batch-job-28qf4 1/1 Running 0 4s
```

After the batch job is complete

```
1 kubectl get po
2 NAME READY STATUS RESTARTS AGE
3 batch-job-k4trj 0/1 Completed 0 110s
```

Check the logs

```
1 kubectl logs batch-job-k4trj
```

```
2 Starting the Batch Job – Tue Mar 31 2020 05:54:11 GMT+0000 (Coordinated Un
3 Completed the Job – Tue Mar 31 2020 05:54:16 GMT+0000 (Coordinated Univers
```

Delete the Job

```
1 kubectl delete job batch-job
2 job.batch "batch-job" deleted
```

Running multiple instance of the Job

Job definition

```
1 apiVersion: batch/v1
2 kind: Job
3 metadata:
4    name: multi-batch-job
5 spec:
6    completions: 5
7    template:
8    metadata:
9    labels:
10     app: batch-job
11    spec:
12    restartPolicy: OnFailure
13    containers:
14    - name: job
15    image: classpathio/batch-job
```

Creating the Job definition

```
1 kubectl create -f hn-multi-batch-job.yaml
2 job.batch/multi-batch-job created
```

Test the Job

```
    kubectl get jobs
    NAME COMPLETIONS DURATION AGE
    multi-batch-job 1/5 18s 18s
```

Verify the Pods

```
kubectl get po
  NAME
                         READY
                                 STATUS
                                            RESTARTS
                                                      AGE
                         0/1
3 multi-batch-job-cmcsn
                                 Completed 0
                                                      91s
                         0/1
0/1
0/1
4 multi-batch-job-gxwsg
                                 Completed 0
                                                      72s
5 multi-batch-job-rmvnb
                                 Completed 0
                                                      100s
                                 Completed 0
6 multi-batch-job-sjzlb
                                                      82s
7 multi-batch-job-zgtvh
                         0/1
                                 Completed 0
                                                      109s
```

Deleting multi-batch-job resource

```
1 kubectl delete job multi-batch-job
2 job.batch "multi-batch-job" deleted
```

Running multiple instance of Job in parallel

```
1 apiVersion: batch/v1
2 kind: Job
3 metadata:
4    name: multi-batch-job
5 spec:
6    completions: 5
7    parallelism: 2
8    template:
9    metadata:
10    labels:
11    app: batch-job
12    spec:
13    restartPolicy: OnFailure
14    containers:
15    - name: job
16    image: classpathio/batch-job
```

Create the job resource

- kubectl create -f hn-multi-batch-job-paralle l.yaml
- job.batch/multi-batch-job created

```
    kubectl get jobs
    NAME COMPLETIONS DURATION AGE
    multi-batch-job 0/5 7s 7s
```

Check the status of pods

```
kubectl get po
  NAME
                          READY
                                  STATUS
                                              RESTARTS
                                                         AGE
  multi-batch-job-drb98
                          0/1
                                  Completed
                                              0
                                                         15s
4 multi-batch-job-hvb6m
                          1/1
                                  Running
                                              0
                                                         4s
5 multi-batch-job-m6txz
                          1/1
                                  Running
                                              0
                                                         6s
6 multi-batch-job-vrmnr
                          0/1
                                  Completed
                                              0
                                                         15s
                          0/1
  multi-batch-job-vxr6j
                                  Completed
                                                         15s
```

After some time, check the status of pods

```
kubectl get po
  NAME
                          READY
                                  STATUS
                                              RESTARTS
                                                         AGE
                          0/1
3 multi-batch-job-drb98
                                  Completed
                                              0
                                                         87s
                          0/1
4 multi-batch-job-hvb6m
                                  Completed
                                                         76s
                                              0
5 multi-batch-job-m6txz
                          0/1
                                  Completed
                                              0
                                                         78s
  multi-batch-job-vrmnr
                          0/1
                                  Completed
                                                         87s
                                              0
  multi-batch-job-vxr6j
                          0/1
                                  Completed
                                              0
                                                         87s
```

Delete the multi-batch-job resource

```
1 kubectl delete job multi-batch-job
2 job.batch "multi-batch-job" deleted
```

CronJob Resource

Creating CronJob Resource

```
apiVersion: batch/v1beta1
2 kind: CronJob
  metadata:
   name: batch-job-every-fifteen-minutes
5 spec:
    schedule: "0,15,30,45 * * * *"
    startingDeadlineSeconds: 15
    jobTemplate:
      spec:
        template:
          metadata:
            labels:
              app: periodic-batch-job
            restartPolicy: OnFailure
            containers:
             - name: main
               image: luksa/batch-job
```

Creating a CronJob from the definition

```
1 kubectl create -f hn-cron-job.yaml
2 cronjob.batch/batch-job-every-fifteen-minutes created
```

Fetching the cronjobs

```
1 kubectl get cronjobs
```

```
2 NAME SUSPEND ACTIVE
3 batch-job-every-fifteen-minutes 0,15,30,45 * * * * False 0
4 95s
```

Deleting cron job resource

- kubectl delete cronjob batch-job-every-fifte en-minutes
- cronjob.batch "batch-job-every-fifteen-minutes" deleted

Services

Services

- K8s resource
- Single point of entry to a group of pods
- · Has a constant IP address and port
- Provides Load balancing capabilities and allow to add/remove pods to/from clusters
- Label selector is used to group pods belongining to the service

Creating a Service resource

Creating a Service using the YAML definition

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4    name: hostname-service
5 spec:
6    ports:
7    - port: 80
8        targetPort: 8111
9    selector:
10    app: hn-service
```

Creating the service resource

```
1 kubectl create -f hn-service.yaml
2 service/hostname-service created
```

Checking the service

kubectl get svc					
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
hostname-service	ClusterIP	100.69.124.66	<none></none>	80/TCP	72s
kubernetes	ClusterIP	100.64.0.1	<none></none>	443/TCP	18m

- Once the Service is created, the pods can be accessed via the cluster IP
- The clusterIP can be used by all the other pods in the cluster
- The clusterIP is visible only within the cluster

Different ways to access the service from within the cluster

- Create a Pod and then invoke the services from within the pod
- ssh into the node and use the curl command
- use the kubectl exec command to issue a curl request from withing a pod

Testing the service from a pod

- Run the kubectl get pods to fetch the pods
- Obtaing the cluster IP using the kubetcl get sv command
- Call the service from the target pod using

Create the Pods using the RS which has the label as hn-service

- 1 kubectl create -f hn-service-rs.yaml
- 2 replicaset.apps/hn-service-rs created

List all the Pods

```
kubectl get po -L app
  NAME
                        READY
                                STATUS
                                          RESTARTS
                                                     AGE
                                                          APP
  hn-service-rs-6h7rf
                        1/1
                                Running
                                          0
                                                     63s
                                                          hn-service
4 hn-service-rs-9qrgq
                        1/1
                                Running
                                                     63s
                                                          hn-service
                                          0
  hn-service-rs-b9r88
                        1/1
                                Running
                                          0
                                                     63s
                                                          hn-service
```

```
kubectl exec hn-service-rs-w2llk -- curl -s http://100.69.124.66
{ "hostname":"hn-service-rs-r5kzs" }

kubectl exec hn-service-rs-w2llk -- curl -s http://100.69.124.66
{ "hostname":"hn-service-rs-kvss9" }

kubectl exec hn-service-rs-w2llk -- curl -s http://100.69.124.66
{ "hostname":"hn-service-rs-w2llk -- curl -s http://100.69.124.66
```

inside the pod.

Configuring session affinity

If we want all the requests to go to a the same pod, then use the sessionAffinity to ClientIP (None is the default)

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4    name: hostname-service-affinity
5 spec:
6    sessionAffinity: ClientIP
7    ports:
8    - port: 80
9     targetPort: 8111
10    selector:
11    app: hn-service
```

! Note: There is no Cookie based session affinity since K8s does not operate at HTTP level. Services deal with TCP and UDP packets and dont care about the payload that they carry.

Create a Service with Client affinity

```
kubectl create -f hn-service-client-affinity.yaml
```

List the services

```
1 kubectl get svc
                            TYPE
                                        CLUSTER-IP
                                                                    PORT
                                                       EXTERNAL-IP
                            ClusterIP
                                       100.69.124.66
3 hostname-service
                                                                    80/T
                                                       <none>
4 hostname-service-affinity ClusterIP
                                       100.70.128.16
                                                                    80/T
                                                       <none>
  kubernetes
                            ClusterIP
                                        100.64.0.1
                                                                    443/
                                                       <none>
```

Affinity will always send the request to the same pod

```
kubectl exec hn-service-rs-w2llk -- curl -s http://100.70.128.16
{ "hostname":"hn-service-rs-kvss9" }

kubectl exec hn-service-rs-w2llk -- curl -s http://100.70.128.16
{ "hostname":"hn-service-rs-kvss9" }

kubectl exec hn-service-rs-w2llk -- curl -s http://100.70.128.16
{ "hostname":"hn-service-rs-kvss9" }

kubectl exec hn-service-rs-kvss9" }

kubectl exec hn-service-rs-w2llk -- curl -s http://100.70.128.16
{ "hostname":"hn-service-rs-kvss9" }
```

Exposing multiple ports in the same Service

A single service can expose multiple port to forward to the pod's port

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4    name: hostname-service
5 spec:
6    ports:
7    - name: http
8         port: 80
9         targetPort: 8080
10    - name: https
11         port: 443
12         targetPort: 8443
13         selector:
14         app: hn-service
```

Using named ports

Target port can also be referred by name instead of port numbers. Name can be given for pod's port

First in the pod definition, define the name and the container port

```
kind: ReplicaSet
metadata:
name: hn-service-named-port-rs
 replicas: 3
selector:
    matchLabels:
      appNname: hn-service
template:
   metadata:
     labels:
        appNname: hn-service
        app: hn-service-named-port
   spec:
     containers:
      - name: hn-service-container
        ports:
        - name: http
           containerPort: 8111
        image: classpathio/hn-service
```

In the service definition

```
1 apiVersion: v1
2 kind: Service
3 metadata:
   name: hostname-service-named-port
5 spec:
   ports:
   - port: 80
     targetPort: http
   selector:
     app: hn-service-named-port
```



Advantages - Enables to change port numbers without having to change the service spec

Discovering Services

Through environmental variables When a pod starts, K8s initializes a set of environmental variables pointing to each service at that moment. These environmental variables are set only if the service is created before creating the pod.

Excercise

1. Delete all the pods by deleting the ReplicaSet

```
kubectl delete rs hn-service-named-port-rs
replicaset.apps "hn-service-named-port-rs" deleted
```

1. list all the pods after creating the Replica-Set

```
kubectl create -f hn-service-rs.yaml
```

List all the Pods

```
kubectl get po --show-labels
                        READY
                                                          LABELS
  NAME
                                STATUS
                                         RESTARTS
                                                    AGE
  hn-service-rs-fvhkw
                        1/1
                                Running
                                                    45s
                                                          app=hn-service
                                         0
4 hn-service-rs-mqmmc
                        1/1
                                Running
                                         0
                                                    45s
                                                          app=hn-service
  hn-service-rs-wbmg8
                        1/1
                                Running
                                         0
                                                    45s
                                                          app=hn-service
```

1. List the environmentsal variables by running the env command on the target pod

```
kubectl exec hn-service-rs-mqmmc env
```

Output

```
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/usr/lib
   HOSTNAME=hn-service-rs-mqmmc
   KUBERNETES_SERVICE_HOST=100.64.0.1
4 HOSTNAME_SERVICE_NAMED_PORT_PORT_80_TCP_ADDR=100.66.116.116
   KUBERNETES_SERVICE_PORT=443
   KUBERNETES_PORT=tcp://100.64.0.1:443
   HOSTNAME_SERVICE_NAMED_PORT_PORT=tcp://100.66.116.116:80
   HOSTNAME_SERVICE_NAMED_PORT_PORT_80_TCP=tcp://100.66.116.116:80
   KUBERNETES_SERVICE_PORT_HTTPS=443
   KUBERNETES_PORT_443_TCP_PROT0=tcp
11 KUBERNETES_PORT_443_TCP_PORT=443
   KUBERNETES_PORT_443_TCP_ADDR=100.64.0.1
   HOSTNAME_SERVICE_NAMED_PORT_SERVICE_HOST=100.66.116.116
   HOSTNAME_SERVICE_NAMED_PORT_SERVICE_PORT=80
15 HOSTNAME_SERVICE_NAMED_PORT_PORT_80_TCP_PROTO=tcp
   HOSTNAME_SERVICE_NAMED_PORT_PORT_80_TCP_PORT=80
   KUBERNETES_PORT_443_TCP=tcp://100.64.0.1:443
18 LANG=C.UTF-8
19 JAVA_HOME=/usr/lib/jvm/java-1.8-openjdk
   JAVA_VERSION=8u212
   JAVA_ALPINE_VERSION=8.212.04-r0
22 HOME=/home/spring
```

When a front-end pod, wants to communicate with the backend-pod, the backend-pod should expose a service which will be used by the front-end pod

(i) Note - Dashes in the service-name will be converted to _ in the environmental variables and all letters will be converted to UPPERCASE.

Discovering services through DNS

- The kube-system service contains a pod named kube-dns .
- All the pods in the cluster are automatically configured to use.
- K8s modified each containers DNS entry in /etc/resolv.conf file
- Any DNS query performed by the process running inside the pod will be handled by K8s DNS server which knows about all the services running inside the system.
- Each service gets a DNS entry in the internal DNS server
- Client pods that know the name of the service, can access the service using the FQDN instead of resorting to the environmental variables

Example: if the service name is backend-database, then the service name will be

hostname-service.default.svc.cluster.local

The svc.cluster.local is optional and the service name can be reffered by backend-database omiitting the suffix.

Excercise:

kubectl exec -it hn-service-rs-mcmp9 bash

Now, inside the container, run the curl command

```
bash-4.4$ curl http://hostname-service.default.svc.cluster.local
{ "hostname":"hn-service-rs-mcmp9" }

bash-4.4$ curl http://hostname-service.default
{ "hostname":"hn-service-rs-z25j6" }

bash-4.4$ curl http://hostname-service
{ "hostname":"hn-service-rs-mcmp9" }

bash-4.4$ cat /etc/resolv.conf
nameserver 100.64.0.10
search default.svc.cluster.local svc.cluster.local cluster.local ap-south-options ndots:5
```

Note: You cannot ping the service because, the service's cluster IP is a virtual IP and only has meaning when combined with the service port.

Connecting to Services outside the cluster

- Client pods can connect to external services outside the cluster
- To leverage load balancing
- Service discovery

Using Service endpoints

Endpoint resource is the link between service and the pods.

```
kubectl describe svc hostname-service
                     hostname-service
3 Namespace:
                     default
4 Labels:
                     <none>
5 Annotations:
                     <none>
6 Selector:
                     app=hn-service
7 Type:
                     ClusterIP
8 IP:
                     100.67.163.196
                     <unset> 80/TCP
9 Port:
10 TargetPort:
                     8111/TCP
11 Endpoints:
                     100.96.1.5:8111,100.96.2.6:8111,100.96.3.4:8111
12 Session Affinity:
                     None
13 Events:
                      <none>
```

- Endpoint resource is a list of IP addresses and ports exposing a service
- You can fetch the Endpoint resource just like any other resource

```
kubectl get endpoints hostname-service
NAME ENDPOINTS AGE
hostname-service 100.96.1.5:8111,100.96.2.6:8111,100.96.3.4:8111 45m
```

Manually configuring the Service endpoints

- · Can be updated manually
- If we create a service without pod selector, endpoints will not even be created
- We then need to create the Endpoint resource manually to specify the list of endpoints for the service

Creating a service without a selector

- Before continuing the below steps delete the Replica-Set, Service and Endpoint resource
- 1. Create a service without the pod selector

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: hn-endpiont-service
5 spec:
6   ports:
7  - port: 80
```

Create the Service from the manifest file

kubectl create -f hn-service-without-endpoints.yaml

Creating a Endpoint resource for a service without a selector

```
1 kubectl create -f hn-endpoint.yaml
2 endpoints/hn-endpiont-service created
```

Create a ReplicaSet

```
1 kubectl create -f hn-service-rs .yaml
2 replicaset.apps/hn-service-rs created
```

List the Pods

```
kubectl get po
                        READY
                                STATUS
                                          RESTARTS
  NAME
                                                     AGE
                        1/\overline{1}
3 hn-service-rs-6lwvs
                                Running
                                          0
                                                     43s
4 hn-service-rs-984s4
                        1/1
                                Running
                                          0
                                                     43s
5 hn-service-rs-n46bk
                        1/1
                                Running
                                          0
                                                     43s
```

- The name of the service should match the service name
- To migrate the external service to pods running inside K8s, add the selector to the service, thereby makings its endpoints managed automatically
- Also, by removing the selector, from the service, the K8s stops updating its endpoints.

Alias for external service

• Specify the service with type field set to ExternalName

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4    name: external-service
5 spec:
6    type: ExternalName
7    externalName: someapi.somecompany.com
8    ports:
9    - port: 80
```

Exposing services to external Clients

Three ways to make the service accessible to extrenal clients

- Setting the service type to NodePort
 - Opens the port at the cluster node level
- Setting the service type to LoadBalancer
 - Extension of NodePort
 - Makes services accessible thourng a dedicated Load Balancer
 - Redirects traffic to the node port across all the nodes
 - Clients connect to the service using the Load balancer IP address
- Creating the service as Ingress resource
 - Operates at the HTTP level

NodePort service

- All the nodes will have a dedicated port opened at the node level
- Can be accessed through Service's internal cluster IP and reserverd node port

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4 name: kubia-nodeport
5 spec:
6 type: NodePort
7 ports:
8 - port: 80
9 targetPort: 8080
10 nodePort: 30123
11 selector:
12 app: kubia
```

List the service

```
1 $ kubectl get svc kubia-nodeport
2 NAME CLUSTER-IP EXTERNAL-IP PORT(S) AGE
3 kubia-nodeport 10.111.254.223 <nodes> 80:30123/TCP 2m
```

i Note: The external ip now shows as <nodes>

Invoking the service using the nodes

```
1 $ curl http://130.211.97.55:30123
2 You've hit kubia-ym8or
3 $ curl http://130.211.99.206:30123
4 You've hit kubia-xueq1
```

i Need for the Load Balancer: The Client's cannot access the service in case of the Node failure

External Load Balancer

- Extension of NodePort
- Featuer of cloud providers
- K8s cluster supports automatic provision of load balancer from the cloud infrastructure
- LB will havve public IP and client's will connect to the LB's IP address
- If the cloud provider does not support LB, then the Nodeport will be used
- Service will behave like a NodePort

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4    name: hn-service-loadbalancer
5 spec:
6    type: LoadBalancer
7    ports:
8    - port: 80
9     targetPort: 8111
10    selector:
11    app: hn-service
```

```
kubectl create -f hn-service-loadbalancer.yamlservice/hn-service-loadbalancer created
```

```
kubectl get svc
                                       CLUSTER-IP
NAME
                        TYPE
                                                      EXTERNAL-IP
                                           PORT(S)
                                                         AGE
hn-service-loadbalancer
                        LoadBalancer
                                       100.67.50.236
                                                      aa8b726f0507d43dd
kubernetes
                        ClusterIP
                                       100.64.0.1
                                                      <none>
                                           443/TCP
                                                         15m
```

Bring up the pods with ReplicaSet

```
1 kubectl get po --show-labels
2 NAME READY STATUS RESTARTS AGE LABELS
```

```
1/1
                                        0
hn-service-rs-9pk89
                              Running
                                                   31s
                                                          app=hn-service
hn-service-rs-np87q
                      1/1
                              Running
                                        0
                                                   31s
                                                          app=hn-service
hn-service-rs-tv6xq
                      1/1
                              Running
                                        0
                                                   31s
                                                          app=hn-service
```

```
1 curl http://aa8b726f0507d43dda45a79f5ebe881d -722734132.ap-south-1.elb.ama
2 { "hostname":"hn-service-rs-9pk89" }
```

To reduce the additional network hop, set the external Traffic Policy inside the spec section and set the value to Local

```
1 spec:
2 externalTrafficPolicy: Local
3 ...
```

! Note:

- By using this feature, if the pod does not exists, then the connection will hang
- Also, the load balancer will not the connection will be forwarded to the same pod

Ingress resource

- Each LoadBalancer service requires its own load balancer with its own Pulic IP address
- Single Ingress requires only on public IP address when providing access to multiple services
- The host and the path will determin the service to be forwarded the request to.
- Operates at the HTTP layer and hence provide features like cookie based session affinity

Creating Ingress resource

Creating an Ingress resource

```
1 apiVersion: extensions/v1beta1
2 kind: Ingress
3 metadata:
4    name: hn-service-ingress
5 spec:
6    rules:
7    - host: hostnameservice.classpath.com
8    http:
9    paths:
10    - path: /
11    backend:
12    serviceName: hostname-service
13    servicePort: 80
```

```
kubectl create -f hn-service-ingress.yamlingress.extensions/hn-service-ingress created
```

```
    kubectl get ingress
    NAME HOSTS AGE
    hn-service-ingress hostnameservice.classpath.com 80 21s
```

Create the Service

```
1 kubectl create -f hn-service.yaml
2 service/hostname-service created
```

Create the ReplicaSet backed with the Servcie created above

```
1 create -f hn-service-rs.yaml
```

replicaset.apps/hn-service-rs created

List the Service

kubectl get svc					
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	
hostname-service	ClusterIP	100.65.248.53	<none></none>	80/TCP	
kubernetes	ClusterIP	100.64.0.1	<none></none>	443/TCP	

List the Pods

```
kubectl get pods --show-labels
  NAME
                       READY
                              STATUS
                                       RESTARTS
                                                  AGE
                                                       LABELS
3 hn-service-rs-dzgml 1/1
                                                       app=hn-service
                              Running
                                                  33s
4 hn-service-rs-nll4z
                       1/1
                              Running
                                                  33s
                                                       app=hn-service
                                       0
5 hn-service-rs-q9g7g 1/1
                              Running
                                                  33s
                                                       app=hn-service
                                        0
```

List the ingress-service

Perform nslookup to fetch the IP address

```
nslookup a9a7caaf3056d4566aa316dffabeb368-17 8615173.ap-south-1.elb.amazon
Server: 172.31.0.2
Address: 172.31.0.2#53

Non-authoritative answer:
Name: a9a7caaf3056d4566aa316dffabeb368-178615173.ap-south-1.elb.amazonaw
Address: 13.232.18.133
Name: a9a7caaf3056d4566aa316dffabeb368-178615173.ap-south-1.elb.amazonaw
```

```
9 Address: 13.126.101.254
```

Configure the DNS to point the domain to the IP address in the /etc/hosts entry file

```
13.126.101.254 hostnameservice.classpath.com
```

Accessing the service with the domain name

```
1 curl http://hostnameservice.classpath.com
2 { "hostname":"hn-service-rs-q9g7g" }
```

Exposing multiple services through the same Ingress

Expose multiple paths on the same host

```
1 apiVersion: extensions/v1beta1
2 kind: Ingress
3 metadata:
   name: hn-service-ingress
5 spec:
    rules:
   host: hostnameservice.classpath.com
     http:
        paths:
        - path: /
          backend:
            serviceName: hostname-service
            servicePort: 80
       - path: /hosts
         backend:
           serviceName: hostname-service
           servicePort: 80
```

Expose multiple hosts

```
spec:
 rules:
 host: hostnameservice.classpath.com
    http:
     paths:
     - path: /
       backend:
          serviceName: hostname-service
          servicePort: 80
- host: bar.example.com
  http:
     paths:
     - path: /
       backend:
         serviceName: bar
          servicePort: 80
```

Configuring TLS

• Delete the hn-service-ingress resource

Create a secret

```
1 kubectl create -f hn-secret.yaml
2 secret/hn-secret created
```

```
    kubectl get ingresses
    NAME HOSTS AGE
    hn-service-ingress hostnameservice.classpath.com
    80 34m
```

```
kubectl delete ingress hn-service-ingress
ingress.extensions "hn-service-ingress" deleted
```

· Create a Manifest file

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: hn-service-ingress-https
spec:
 tls:
 - hosts:
    hostnameservice.classpath.com
    secretName: hn-secret
  rules:
 - host: hostnameservice.classpath.com
      paths:
      - path: /
        backend:
          serviceName: hostname-service
          servicePort: 80
```

```
kubectl create -f hn-service-ingress-https.yamlingress.extensions/hn-service-ingress-https created
```

```
    kubectl get ingresses
    NAME
    HOSTS
    hn-service-ingress-https
    hostnameservice.classpath.com
    80, 4
```

Accessing the service

```
1 curl https://hostnameservice.classpath.com --insecure
2 { "hostname":"hn-service-rs-q9g7g" }
```

Readiness Probe

- Invoked periodically and determine whether the specified pod should receive client requests or not.
- If the container's rediness probe returns success, it's signalling that it is ready to accept client's requests.
- Defining readiness probe is specific to each container
- K8s can check if the container is ready by hitting the GET request specified
- Its the container's responsibility to accepts the GET endpoint and perform the checks and return the status.
- When a container is started, K8s can be configured to perform the readiness probe after a configured amount of time
- After the first readiness probe, the K8s performs the probe frequently
- If the pod is not ready, it is removed from the service
- If the pod is ready, it is added to the service
- Unlike the liveness probe, the pod will not be killed or restarted
- Liveness probe is to keep healthy pods by killing and replacing with the healthy pods
- Readiness probe is to ensure that only the pods which can accept client's request are added to the service

Types of Readiness Probes

- Exec Process is executed. Process exit status code determines the readiness
- HTTP GET HTTP status code determines the readiness of the container
- TCP Socket opens up the TCP connection to a specified port. The connection establishment determines, if the service is ready.

Creating a Readiness probe

Create the ReplicaSet definition

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
name: hn-service-readinessrs
spec:
replicas: 3
selector:
matchLabels:
```

```
papp: nodejs-readiness
template:
metadata:
labels:
app: nodejs-readiness
spec:
containers:
name: nodejs-readiness-container
image: classpathio/nodejs-readiness-app
readinessProbe:
httpGet:
path: /readiness
port: 8080
initialDelaySeconds: 10
periodSeconds: 5
```

```
1 kubectl create -f hn-service-readiness-rs.yaml
2 replicaset.apps/hn-service-readinessrs created
```

- If the file exists, it returns 0 indicating success
- Non 0 if the file does not exists indicating failure

List all the pods

```
kubectl get po
NAME
                             READY
                                     STATUS
                                              RESTARTS
                                                         AGE
hn-service-readinessrs-52m69
                                     Running
                             0/1
                                                         71s
hn-service-readinessrs-87mk4
                             0/1
                                     Running 0
                                                         71s
hn-service-readinessrs-rbzcx 0/1
                                     Running
                                             0
                                                         71s
```

```
1 kubectl exec -it hn-service-readinessrs-87mk 4 bash
2 root@hn-service-readinessrs-87mk4:/usr/src/app#
```

```
curl http://localhost:8080/readiness
{"status":"Not up"}
```

```
4 curl http://localhost:8080/init -XPOST
5 {"status":"success"}
```

```
kubectl get po
NAME
                                READY
                                        STATUS
                                                  RESTARTS
                                                              AGE
hn-service-readinessrs-52m69
                                0/1
                                        Running
                                                              4m26s
hn-service-readinessrs-87mk4
                                1/1
                                                              4m26s
                                        Running
                                                  0
hn-service-readinessrs-rbzcx
                                                              4m26s
                                0/1
                                        Running
```

Hit the Service from the ec2-instance

```
kubectl exec hn-service-readinessrs-87mk4 cu rl http://localhost:8080/rea
 % Total % Received % Xferd Average Speed
                                              Time
                                                     Time
                                                              Time Cu
                               Dload Upload
                                                              Left Sp
                                              Total
                                                     Spent
100
      15 100
                 15
                      0
                            0
                                2035
                                         0 --:--:-- --:--:--
{"status":"Up"}
```

Optional steps

· Make the Pod not-ready again

All Pods should now be in Not-Ready state again

```
kubectl get po
NAME
                                READY
                                        STATUS
                                                  RESTARTS
                                                             AGE
hn-service-readinessrs-52m69
                                0/1
                                        Running
                                                              8m30s
hn-service-readinessrs-87mk4
                                                              8m30s
                                0/1
                                        Running
hn-service-readinessrs-rbzcx
                                0/1
                                        Running
                                                              8m30s
```

Headless service

- In cases when the clients need to talk to all the pods
- Can get the IP address of all the pods by connecting to the API server, but not ideal since it will be decoupled with the K8s.
- Setting up the ClusterIP to None, K8s does not return the cluster IP. The DNS server will return the list of pod IP addresses.
- Clients talk directly to the pod's IP address rather than the service proxy

Creating a headless service

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4    name: hostname-service
5 spec:
6    ClusterIP: None
7    ports:
8    - name: http
9    port: 80
10    targetPort: 8080
11    - name: https
12    port: 443
13    targetPort: 8443
14    selector:
15    app: hn-service
```

Create the service using the kubectl create command

Inspect the Cluser IP field by running the kubectl describe command

```
$ kubectl exec <pod name>
```

Volumes

Volume

- · Process running inside the pods share the resources like CPU, RAM, Network interfaces
- · Disks are not shared, container inside the pod has its own isolated filesystem
- File system comes from the container's image
- Volumes are not top level resources like RC , RS
- Is part of the pod definition
- Share the same lifecycle as of the pod
- Volume's content will be persisted across pod restarts
- · Volumne is available for all the containers inside the pod
- It must be mounted in each container that need access to the volume
- Can mount the volume in any location of the file system

Available Volume types

- emptyDir a simple directory
- hostPath a directory on the node mounted to the pods
- gitRepo Volume initialized by cloning the contents from the git repository
- nfs NFS share mount into the pod
- awsElasticBlockStorage AWS EBS volume
- azureDisk Azure
- persistentVolumeClaim dynamically provisioned persistent storage

Using emptyDir volume

- Simple solution
- · Starts with a empty directory
- Volume lifetime is tied to the pod's lifetime
- Useful for sharing data between containers in the same pod

```
apiVersion: v1
kind: Pod
metadata:
  name: fortune
spec:
 containers:
  - image: luksa/fortune
    name: html-generator
    volumeMounts:
   - name: html
      mountPath: /var/htdocs
 - image: nginx:alpine
    name: web-server
   volumeMounts:
    - name: html
      mountPath: /usr/share/nginx/html
      readOnly: true
    ports:
    - containerPort: 80
      protocol: TCP
 volumes:
 - name: html
    emptyDir: {}
```

```
    kubectl get po
    NAME READY STATUS RESTARTS AGE
    fortune 2/2 Running 0 15s
```

Testing the pod

```
1 $ kubectl port-forward fortune 8080:80
2 Forwarding from 127.0.0.1:8080 -> 80
3 Forwarding from [::1]:8080 -> 80
```

Test the pod by invoking the HTTP request

```
    curl http://localhost:8080
    Noise proves nothing. Often a hen who has merely laid an egg cackles
    as if she laid an asteroid.
```

Try invoking the HTTP request after some time

```
    curl http://localhost:8080
    Q: What's the difference between a Mac and an Etch-a-Sketch?
    A: You don't have to shake the Mac to clear the screen.
```

- You should see a different message this time
- The volume is created on the worker node hosting the pod
- You can also set the medium to Memory which creates an in-memory

Using Git Repo

- Similar to emptyDir but the content will be populated by cloning from git repository
- It is not kept in-sync in case of updates to the repository

```
apiVersion: v1
2 kind: Pod
3 metadata:
   name: gitrepo-volume-pod
5 spec:
   containers:
   - image: nginx:alpine
     name: web-server
     volumeMounts:
     - name: html
       mountPath: /usr/share/nginx/html
       readOnly: true
      ports:
     - containerPort: 80
        protocol: TCP
   volumes:
    - name: html
      gitRepo:
        repository: https://github.com/luksa/kubia-website-example.git
```

```
revision: master
directory: .
```

```
1 kubectl create -f gitrepo-volume-rs.yaml
2 replicaset.apps/gitrepo-rs created
```

```
1 kubectl get rs
2 NAME DESIRED CURRENT READY AGE
3 gitrepo-rs 2 2 27s
```

Two Options

- Using Services and Ingress
- Port Forwarding

Service and Ingress

```
1 kubectl create -f git-repo-service.yaml
2 service/git-repo-service created
```

Create a Ingress Manifest

```
1 apiVersion: extensions/v1beta1
2 kind: Ingress
3 metadata:
4    name: git-repo-ingress
5 spec:
6    tls:
7    - hosts:
8         - blog.classpath.io
9         secretName: cafe-secret
10    rules:
11         - host: blog.classpath.io
12         http:
```

```
paths:
- path: /
- path: /
backend:
- serviceName: git-repo-service
- servicePort: 80
```

```
1 kubectl create -f volumes/git-repo-ingress-https.yaml
2 ingress.extensions/hn-service-ingress created
```

```
    kubectl get ingresses
    NAME HOSTS ADDRESS PORTS A
    hn-service-ingress hostnameservice.classpath.com 80, 443 4
```

Edit the /etc/hosts entry with the hostnameservice.classpath.io

```
13.232.188.216 curl https://blog.classpath.io
```

Port forwarding - Simple :-)

```
1 kubectl port-forward gitrepo-rs-4x4w5 8080:80
2 Forwarding from 127.0.0.1:8080 -> 80
3 Forwarding from [::1]:8080 -> 80
```

In another terminal run the curl command

```
curl http://localhost:8080
chtml>
description
formula to the state of the stat
```

Delete the Pod

```
1 kubectl delete po gitrepo-rs-4x4w5
2 pod "gitrepo-rs-4x4w5" deleted
```

List the Pods

```
1 kubectl get po
2 NAME READY STATUS RESTARTS AGE
3 gitrepo-rs-4kr6b 1/1 Running 0 39s
4 gitrepo-rs-m97wr 1/1 Running 0 50m
```

```
kubectl port-forward gitrepo-rs-4kr6b 8080:80
Forwarding from 127.0.0.1:8080 -> 80
Forwarding from [::1]:8080 -> 80
```

Old Pod with older version of git

```
kubectl port-forward gitrepo-rs-m97wr 8080:80
Forwarding from 127.0.0.1:8080 -> 80
Forwarding from [::1]:8080 -> 80
Handling connection for 8080
```

```
curl http://localhost:8080
chtml>
download
height in the state of the state of
```

• When the pod is deleted, the volume is also deleted

Hostpath Volume

- Volume is not deleted even if the pod is deleted
- Volume's content is stored on the node
- System level services (DaemonSet) can leverage this feature
- Makes the pod sensitive to the node it is scheduled on

Using Persistent storage



Create a Volume in Elastic Block Storage and create a tag with key as KubernetesCluster and value as classpath.k8s.local

Also create the Volume inside the same availablity zone

Create a Pod definition

```
apiVersion: v1
2 kind: Pod
3 metadata:
   name: mongodb
5 spec:
6 volumes:
   name: mongodb-data
     awsElasticBlockStore:
        volumeID: vol-09267ef3713fb2ede
        fsType: ext4
   containers:
   - image: mongo
     name: mongodb
     volumeMounts:
     name: mongodb-data
      mountPath: /data/db
     ports:
     - containerPort: 27017
        protocol: TCP
```

- Pod contains a single container and a single volume baked by the Persistent disk
- Mounted the volume at /data/db
- Now run the MongoDB inside the shell

```
kubectl create -f mongo-db-data.yaml
```

```
kubectl exec -it mongodb mongo
MongoDB shell version: 3.2.8
```

```
connecting to: mongodb://127.0.0.1:27017
Welcome to the MongoDB shell.
For interactive help, type "help".
For more comprehensive documentation, see
    http://docs.mongodb.org/
Questions? Try the support group
    http://groups.google.com/group/mongodb-user
...
```

Insert some documents inside the MongoDB

```
1 > use store
2 switched to db store
3 > db.cart.insert({name:'classpath'})
4 WriteResult({ "nInserted" : 1 })
5 > db.cart.find()
6 { "_id" : ObjectId("5e85f350f5b6271cca763c1d"), "name" : "classpath" }
```

Recreating the pod by deleting the pod

```
kubectl delete pod mongodb
pod "mongodb" deleted

kubectl create -f mongo-db-data.yaml
pod "mongodb" created
```

Retriving the persisted data after recreating the pod

```
kubectl exec -it mongodb mongo

vuse store
switched to db store
blue db.cart.find()
{ "_id" : ObjectId("5e85f350f5b6271cca763c1d"), "name" : "classpath" }
```

Decoupling the pods from underlying storage technology

- K8s allows us to decouple the underlying storage technology
- Makes our application to switch between technology easily

PersistentVolumes and PersistentVolumeClaim

Steps to use the PersistentVolume and PersistentVolumeClaim

Cluster Admin sets up the PersistentVolume

- Cluster admin sets up the storage and registers with K8s
- Registering the volume is done by creating a PersistentVolume resource through K8s
 API Server
- When creating the PersistentVolume, the admin specifies the size and access mode it supports (read / write)

Cluster user then creates a PersistentVolumeClaim

- User creates a PersistentVolumeClaim resource
- Specifies the minimum size and the access mode
- User submits the PersistentVolumeClaim to K8s's API server
- K8s finds the suitable PersistentVolume an binds to the volume to the PersistentVolumeClaim
- The PersistentVolumeClaim is now ready to be mounted on any pods
- Other users cannot use the same persistentVolumeclaim until the resource is release by deleting the resource.

Creating the PersistentVolume

Creating PersistentVolume

```
1 apiVersion: v1
2 kind: PersistentVolume
3 metadata:
4    name: mongodb-pv
5 spec:
```

```
capacity:
storage: 10Gi
sccessModes:
- ReadWriteOnce
- ReadOnlyMany
persistentVolumeReclaimPolicy: Retain
awsElasticBlockStore:
volumeID: vol-09267ef3713fb2ede
fsType: ext4
```

```
kubectl create -f aws-persistent-volume.yaml
```

Fetch the PersistentVolume

```
kubectl get pv
```

Creaing a PersistentVolumeClaim

```
1 apiVersion: v1
2 kind: PersistentVolumeClaim
3 metadata:
4    name: mongodb-pvc
5 spec:
6    resources:
7    requests:
8        storage: 1Gi
9    accessModes:
10    - ReadWriteOnce
11    storageClassName: ""
```

```
kubectl create -f mongodb-pvc.yaml
```

- K8s will find the appropriate PersistentVolume and binds to the claim
- The PersistentVolume should be larger than the PersistentVolumeClaim
- The volume's access modes must include the access modes requested by the claim.

List the PersistentVolumeClaim

```
1 $ kubectl get pvc
2 NAME STATUS VOLUME CAPACITY ACCESSMODES AGE
3 mongodb-pvc Bound mongodb-pv 1Gi RWO,ROX 3s
```

- RWO—ReadWriteOnce— Only a single node can mount the volume for reading and writing.
- ROX—ReadOnlyMany—Multiple nodes can mount the volume for reading.
- RWX—ReadWriteMany—Multiple nodes can mount the volume for both reading and writing.

List the PersistentVolume

```
1 $ kubectl get pv
2 NAME CAPACITY ACCESSMODES STATUS CLAIM AGE
```

Using the PersistentVolumeClaim inside the Pod

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4    name: mongodb
5 spec:
6    containers:
7    - image: mongo
8    name: mongodb
9    volumeMounts:
10    - name: mongodb-data
11    mountPath: /data/db
12    ports:
13    - containerPort: 27017
```

```
protocol: TCP
volumes:
- name: mongodb-data
persistentVolumeClaim:
claimName: mongodb-pvc
```

```
kubectl create -f mongodb-with-pvc.yaml
```

Retrieve the data persisted

```
kubectl exec -it mongodb mongo

vuse store
switched to db store
by db.cart.find()
{ "_id" : ObjectId("5e85f350f5b6271cca763c1d"), "name" : "classpath" }
```

Recycling PersistentVolumeClaim

```
$ kubectl delete pod mongodb
pod "mongodb" deleted
$ kubectl delete pvc mongodb-pvc
persistentvolumeclaim "mongodb-pvc" deleted
```

Status of PersistentVolumeClaim after deleting

```
1 kubectl get pv
2 NAME CAPACITY ACCESSMODES STATUS CLAIM REASON AG
3 mongodb-pv 1Gi RWO,ROX Released default/mongodb-pvc 5m
```

```
1 kubectl get pvc
```

2 NAME STATUS VOLUME CAPACITY ACCESSMODES AGE 3 mongodb-pvc Pending 13s

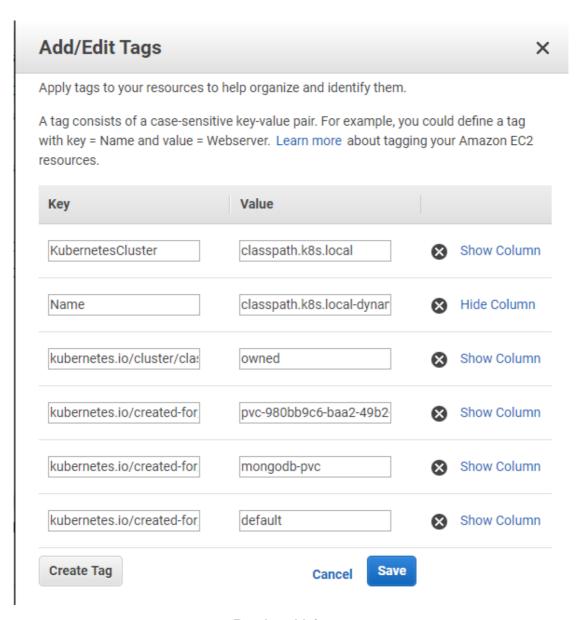
Dynamic provision of PersitentVolumes

- PersistentVolume is still vendor dependent
- Cluster admin will have to configure the space upfront
- Cluster admin can instead deploy a PersistentVolume provisioner
- Define one or more StorageClass in their PersistentVolumeClaim
- Provisioner will take into account when provisioning the persistent storage
- This ways, the provisioner can provision from logically infinite storage and not upfront create the volume.
- StorageClass resource aren't namespace
- K8s have provisioner for most cloud providers
- Custom provisioner should be deployed in case of a on-premise
- Cluster admin will pre-provision one or more StorageClass instead of PersistentVolume

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
    name: standard
provisioner: kubernetes.io/aws-ebs
parameters:
    type: gp2
reclaimPolicy: Retain
allowVolumeExpansion: true
mountOptions:
    - debug
volumeBindingMode: Immediate
```

```
create -f aws-storage-class.yaml
storageclass.storage.k8s.io/standard created
```

2 NAME PROVISIONER AGE
3 default kubernetes.io/aws-ebs 4m47s
4 gp2 (default) kubernetes.io/aws-ebs 4m47s
5 standard kubernetes.io/aws-ebs 70s



PersitentVolume

Creating a PersistentVolumeClaim to use dynamic provisioning.

apiVersion: v1

2 kind: PersistentVolumeClaim

```
metadata:
name: mongodb-sc-pvc
spec:
storageClassName: standard
resources:
requests:
storage: 100Mi
accessModes:
- ReadWriteOnce
```

```
1 kubectl create -f aws-pvc-with-storage-class.yaml
2
3 persistentvolumeclaim/mongodb-sc-pvc created
```

List the PersistentVolumeClaim

```
1 kubectl get pvc
2 NAME STATUS VOLUME CAPAC
3 mongodb-sc-pvc Bound pvc-aa387010-f5fb-4429-8380-628c708e8a26 1Gi
4 standard 19s
```

List the PersistentVolume

```
kubectl get pv
2 NAME
                                            CAPACITY
                                                       ACCESS MODES
                                                                      RECLA
             CLAIM
3 STATUS
                                     STORAGECLASS
                                                    REASON
                                                             AGE
4 pvc-980bb9c6-baa2-49b2-aa84-066862270284
                                                       RWO
                                            1Gi
                                                                      Retai
  Released
             default/mongodb-pvc
                                      standard
                                                             14m
6 pvc-aa387010-f5fb-4429-8380-628c708e8a26
                                           1Gi
                                                       RWO
                                                                      Retai
             default/mongodb-sc-pvc
                                     standard
                                                             71s
  Bound
  pvc-f5346f31-4c44-480b-adb2-fcad0818be46 1Gi
                                                       RWO
                                                                      Retai
                                     standard
  Released default/mongodb-pvc
                                                             3m52s
```

Create the Mongodb Pod

- kubectl create -f mongodb-with-pvc.yaml
- pod/mongodb created

```
Events:
 Type
         Reason
                                 Age
                                       From
       Message
 Normal Scheduled
                                 64s
                                       default-scheduler
       Successfully assigned default/mongodb to ip-172-20-55-45.ap-south-
 Normal SuccessfulAttachVolume 58s
                                       attachdetach-controller
       AttachVolume.Attach succeeded for volume "pvc-aa387010-f5fb-4429-8
 Normal Pulling
                                 55s
                                       kubelet, ip-172-20-55-45.ap-south-
 Normal Pulled
                                       kubelet, ip-172-20-55-45.ap-south-
                                 41s
 Normal Created
                                 41s
                                       kubelet, ip-172-20-55-45.ap-south-
                                       kubelet, ip-172-20-55-45.ap-south-
 Normal Started
                                 41s
```



Assignment:

- 1. Create the Mongodb Pod
- 2. Log into to mongo shell
- 3. Insert the record and exit
- 4. Delete the Pod
- 5. Recreate the Pod
- 6. Log into the Mongo shell
- 7. Find the inserted record

Delete the Pod

- kubectl delete po mongodb
- pod "mongodb" deleted

Delete the PVC

- kubectl delete pvc mongodb-sc-pvc
- persistentvolumeclaim "mongodb-sc-pvc" deleted

Delete the Storage Class

- 1 kubectl delete sc --all
- 2 storageclass.storage.k8s.io "default" deleted
- 3 storageclass.storage.k8s.io "gp2" deleted
- 4 storageclass.storage.k8s.io "standard" deleted

Delete the PV

- 1 kubectl delete pv --all
- persistentvolume "pvc-980bb9c6-baa2-49b2-aa84-066862270284" deleted
- 3 persistentvolume "pvc-aa387010-f5fb-4429-8380-628c708e8a26" deleted
- 4 persistentvolume "pvc-f5346f31-4c44-480b-adb2-fcad0818be46" deleted

ConfigMaps

ConfigMaps

By default when we first start building our application

- First pass the arguments using command line
- Then when the arguments become difficult to manage, move them to external configuration file
- In containerized applications we pass the arguments using the environmental variables. Ex: mysql password

Passing Environmental Variables inside a Container

- The config file has to be baked into the image directly
- Similar to hard coding in the source file since it requires to rebuild the image every time we change the configuration entries
- Anyone having access to the image can see all the information including sensitive information.
- Alternative method would be to mount the volume containing the config file into the container

Ways of Configuring properties

- Passing command line arguments to containers
- Setting custom environmental variable for each container
- Mounting configuration files into containers through a special type of volume

Defining command line arguments in Docker

- Command that gets executed inside the container contains two parts namely command and arguments.
- ENTRYPOINT command to be executed
- CMD specifies the arguments to be passed to the command

Then the image can be run without passing the command

```
docker run <image>
```

Or override the arguments which was set under the CMD in the dockerfile

```
docker run <image> <arguments>
```

```
shell-form - ENTRYPOINT node app.js

exec-form - ENTRYPOINT ["node", "app.js"]
```

i Always use the exec mode of running docker images

Overriding command and arguments in K8s

In K8s we can override both the ENTRYPOINT and the CMD

```
kind: Pod
spec:
containers:
    - image: some/image
command: ["/bin/command"]
args: ["arg1", "arg2", "arg3"]
```

i The command and args fields can't be updated after the pod is created.

-										
I٦		0	v	Or	CO	m	m	9	m	പ
\boldsymbol{L}	u	w	n.	CI.	CU			a		u

Kubernetes Command

ENTRYPOINT	command
CMD	args

Running with custom delay interval

```
apiVersion: v1
kind: Pod
metadata:
name: talkingcow-args
spec:
containers:
- image: classpathio/talking-cow-args
name: talkingcow-args
args: ["4"]
```

To pass more than one or few arguments, you can follow the below syntax

```
1 args:
2 - one
3 - two
4 - "10"
```

Setting Environmental variables for a container

- Containerized applications often use env variables as a source of configuration option
- K8s allows to specify custom list of env variables for each container of the pod
- Currently no option exists to set the env variables at the pod level and inherit by its containers
- The environmental variables cannot be updated after the pod is created

Defining environmental variable in the pod definition

Define env at the container level and not at the pod level

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4    name: talkingcow-env
5 spec:
6    containers:
7    - image: classpathio/talking-cow-env
8    name: talkingcow-env-container
9    env:
10    - name: INTERVAL
11    value: "15"
```

Environment variables can be reffered inside the pod definition

```
1 env:
2 - name: FIRST_VAR
3 value: "foo"
4 - name: SECOND_VAR
5 value: "$(FIRST_VAR)bar"
```

- i The SECOND_VAR value will be "foobar".
- ! Hard coding the **environmental variables** in the pod definition makes it difficult to manage across different environments as we would need to maintain different pod definitions for each environments.

Decouple configurations out of the POD descriptor with the help of ConfigMaps

ConfigMaps

- ConfigMap is a map containing key and value
- Value containing simple strings to full config files
- Application does not need to know that the config map exists
- Application does not read the config map directly
- Contents of the map is either sent as either env variables or as files in a volume
- Can also be passed as command line arguments referred with \${env_var}
- This allows for keeping multiple manifest files with same name, each for different environments

Creating a Config Map

```
kubectl create configmap config-sleep-10 --from-literal=sleep-interval=10
```

(i) All examples use the create command to create the configmaps and secret.

ConfigMap with multiple keys

```
kubectl create configmap myconfigmap
   --from-literal=foo=bar --from-literal=bar=baz --from-literal=one=two
```

View the YAML definition

```
1 apiVersion: v1
   sleep-interval: "4"
4 kind: ConfigMap
5 metadata:
   name: fortune-interval-4
```

Creating the ConfigMap using the yaml definition descriptor

\$ kubectl create -f config-map-interval.yaml

Creating a config entry from the contents of a file

kubectl create configmap my-config --from-file=config-map-files.conf

in the above case, the filename will be used as a key

Can also specify the key explicitly

kubectl create configmap my-config --from-file=customkey=config-file.conf

Import all the files from a directory

kubectl create configmap my-config --from-file=/path/to/dir

in the above case, each file name will be the key with value referring to the contents of the file

Combine all of them like below:

l kubectl create configmap my-config --from-file=foo.json --from-file=bar=f

Getting the values of the configmap inside the container

- Setting an environmental variable
- Command line argument
- Volume to expose ConfigMap entries as files

Passing as Environmental Variable

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4    name: talkingcow-configmap-env
5 spec:
6    containers:
7    - image: classpathio/talking-cow-env
8    name: talkingcow-env-configmap-container
9    env:
10    - name: INTERVAL
11    valueFrom:
12    configMapKeyRef:
13    name: fortune-interval-4
14    key: sleep-interval
```

- interval is the environmenal variable name and the value is picked up from the key named sleep-interval inside the fortune-config ConfigMap
- i If the referenced ConfigMap doesn't exist when you create the pod. Kubernetes schedules the pod normally and tries to run its containers.

The container referencing the non-existing ConfigMap will fail to start, but the other container will start normally.

If you then create the missing ConfigMap, the failed container is started without requiring you to recreate the pod.

Passing all the environmental properties at once

```
1 spec:
2 containers:
3 - image: some-image
4 envFrom:
5 - prefix: CONFIG_
6 configMapRef:
7 name: my-config-map
8 ...
```

Passing as Command Line arguments

- First initialize the env variable with the value of configmap entry
- Refer that variable inside the argument

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4    name: talkingcow-configmap-args
5 spec:
6    containers:
7    - image: classpathio/talking-cow-env
8    name: talkingcow-env-configmap-container
9    env:
10    - name: INTERVAL
11    valueFrom:
12         configMapKeyRef:
13         name: fortune-interval-4
14         key: sleep-interval
15    args: ["$(INTERVAL)"]
```

Exposing configMap entries as volumes

- For passing entire configuration files we can use configMap volume
- Special type of Volume Mount
- Will expose each entry of the ConfigMap as a file
- Process running inside the container can obtain the entry's value by reading the contents of the file
- Can also be used to send simple config entries and single values.

Creating a Configmap Volume

Delete the ConfigMap

```
kubectl delete configmap config-map-name
```

Creating a configmap with the config-map-files stored on the disk

Steps

- 1. Create a new directory called config-map-files
- 2. Store the nginx-conifuguration file inside the directory
- 3. Create a configmap resource with the below command

```
kubectl create configmap nginx-config --from-file=config-map-filesconfigmap/nginx-config created
```

```
kubectl get configmapsNAME DATA AGEnginx-config 2 65s
```

```
kubectl get configmaps nginx-config -o yaml
apiVersion: v1
data:
  nginx-configuration.conf: |
   erver {
     listen
                          82;
     server_name
                        www.classpath.io;
      gzip on;
     gzip_types text/plain application/xml;
      location / {
        root /usr/share/nginx/html;
        index index.html index.htm;
      }
  sleep-interval.conf: |
   4
kind: ConfigMap
metadata:
creationTimestamp: "2020-04-06T04:26:43Z"
name: nginx-config
namespace: default
resourceVersion: "968"
 selfLink: /api/v1/namespaces/default/configmaps/nginx-config
 uid: 51f57c9c-d3d1-4777-8214-c9109f5eec02
```

i For Nginx, the server will load all the configuration file under /etc/nginx/conf.d directory.

We need to place our configuration under that folder and not replace the nginx.conf file

Pod definition

```
apiVersion: v1
2 kind: Pod
3 metadata:
   name: nginx-configmap-volume
5 spec:
   containers:
   - image: nginx:alpine
     name: webserver
     volumeMounts:
     - name: configuration
       mountPath: /etc/nginx/conf.d
       readOnly: true
      ports:
     - containerPort: 80
        protocol: TCP
   volumes:
   - name: configuration
      configMap:
        name: nginx-config
```

Verifying if Nginx is using the Configuration

```
curl -H "Accept-Encoding: gzip" -I http://localhost:8000
HTTP/1.1 200 OK
Server: nginx/1.17.9
Date: Mon, 06 Apr 2020 05:16:10 GMT
Content-Type: text/html
Last-Modified: Tue, 03 Mar 2020 17:36:53 GMT
Connection: keep-alive
ETag: W/"5e5e95b5-264"
Content-Encoding: gzip
```

Examining the contentes of ConfigMap in the Pod

```
    kubectl exec nginx-configmap-volume ls /etc/nginx/conf .d
    nginx-configuration.conf
    sleep-interval
```

- Both entries of ConfigMap have been added as files to the directory
- · Could create multiple configmaps to specific container

Exposing certain ConfigMaps entries in the Volume

To define specific entries as files in ConfigMap volume, use the Volumes items attribute

```
apiVersion: v1
2 kind: Pod
3 metadata:
   name: nginx-configmap-volume
  spec:
   containers:
   - image: nginx:alpine
     name: webserver-config
     volumeMounts:
     - name: configuration
       mountPath: /etc/nginx/conf.d
       readOnly: true
     ports:
     - containerPort: 82
        protocol: TCP
   volumes:
   - name: configuration
     configMap:
        name: nginx-config
        items:
        - key: nginx-configuration.conf
          path: gzip.conf
```

Examining the contentes of ConfigMap in the Pod this time

```
$ $ kubectl exec fortune-configmap-volume -c web-server ls /etc/nginx/conf.d
## my-nginx-config.conf
```

! Mounting a directory hides the existing files in that directory as long as the directory is mounted in Linux systems

Mounting individual ConfigMap entries as files without hiding other files in the directory

• Add the **subPath** property on the **volumeMount** allows to mount either a single file or single directory from the volume instead of mounting the entire volume

```
spec:
containers:
    - image: some/image
volumeMounts:
    - name: myvolume
    mountPath: /etc/someconfig.conf
subPath: myconfig.conf
```

Setting the file permissions for files in the ConfigMap

```
volumes:
volumes:
name: config
configMap:
name: fortune-config
defaultMode: "6600"
```

Updating apps configuration without restarting the app

- Configuration entries passed via env variables and command line arguments cannot be updated while the proces is running
- Using ConfigMap and exposing through volume makes updating the configuration change without restarting the container
- Its up to the process to detect the changes and reload them

Editing the ConfigMap

```
kubectl edit configmap fortune-config
```

Edit the configuration changes

```
$ kubectl exec fortune-configmap-volume -c web-server
2 ⇒ cat /etc/nginx/conf.d/my-nginx-config.conf
```

Signaling the process to reload the configuration

```
$ kubectl exec fortune-configmap-volume -c web-server -- nginx -s reload
```

Now make the request and verty the response

```
kubectl port-forward fortune-configmap-volume 8080:80 &
2 Forwarding from 127.0.0.1:8080 -> 80
3 Forwarding from [::1]:8080 -> 80
4 $ curl -H "Accept-Encoding: gzip" -I localhost:8080
5 HTTP/1.1 200 OK
6 Server: nginx/1.11.1
7 Date: Thu, 18 Aug 2016 11:52:57 GMT
8 Content-Type: text/html
9 Last-Modified: Thu, 18 Aug 2016 11:52:55 GMT
10 Connection: keep-alive
11 ETag: W/"57b5a197-37"
12 Content-Encoding: gzip
```



🔼 If file is mounted instead of whole volume, the file will not be updated

(i) Containers should be immutable

Don't bypass the immutability by updating the configmap while containers are running

If App's do not support dynamic reloading, this is not an issue

Secrets

- Secrets are K8s resource to store sensitive configuration data
- Similar to ConfigMaps Holds key-value pairs
- Can be passed to containers as **environment variables**
- Expose secret entries as files in a Volume
- K8s ensures that the secrets is only distributed to the nodes that run the pods that need the secret
- On the Nodes, the secrets are stored in-memory and not persisted to disk

Default secret token in each container

• Every pod has a secret volume attached to it by default

kubectl descrie pod-name

```
1 Volumes:
2  default-token-cfee9:
3  Type: Secret (a volume populated by a Secret)
4  SecretName: default-token-cfee9
```

Secrets are first class K8s resource

```
kubectl get secrets
```

kubectl describe secrets

Secrets contains three entries

- ca.cert
- namespace
- token

The above represents everything needed to securely talk to API server

```
kubectl exec mypod ls /var/run/secrets/kubernetes.io/serviceaccount/
ca.crt
namespace
token
```

Creating a Secret

```
    openssl genrsa -out https.key 2048
    openssl req -new -x509 -key https.key -out https.cert -days
    3 3650 -subj /CN=www.classpath.io
```

Create a dummy file called domain

```
echo www.classpath.io > domain
```

Create a secret file

kubectl create secret generic nginx-https --from-fil e=certs/https.key --fro

Comparing ConfigMaps and Secrets

- The contents of the Secrets are Base 64 encoded but for ConfigMaps the contents are in plain text
- The secrets can contain binay data and not only plain-text
- Base 64 encoded values allows to store the binary data in YAML or JSON format
- Maximum size of the Secret is 1 MB

```
$ kubectl get secret fortune-https -o yaml
apiVersion: v1
data:
domain: d3d3LmNsYXNzcGF0aC5pbwo=
https.cert: LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JSURCekNDQ...
https.key: LS0tLS1CRUdJTiBSU0EgUFJJVkFURSBLRVktLS0tLQpNSUlFcE...
kind: Secret
metadata:
creationTimestamp: "2020-04-06T07:50:27Z"
name: nginx-https
namespace: default
resourceVersion: "1073"
selfLink: /api/v1/namespaces/default/secrets/nginx-https
uid: 4356b830-3068-43fe-ae50-2963c97bce39
type: Opaque
```

Using String data field

- To set non binay data using the StringData field
- Write only Only used to set values
- When retrieving using the kubectl get -o yaml, the string data will not be shown.
- Will be shown under the data section in base-64 encoded format.

 App need not decode the values but can read the contents or look up from environmental variables and use it directly

```
kind: Secret
priversion: v1
stringData:
key: value
data:
https.cert: LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JSURCekNDQ...
https.key: LS0tLS1CRUdJTiBSU0EgUFJJVkFURSBLRVktLS0tLQpNSUlFcE...
```

Using the Secrets in the pod

Create the configmap with the configuration files

```
kubectl create configmap nginx-conf --from-file=config-maps/config-map-files
```

```
    kubectl get configmaps
    NAME DATA AGE
    nginx-conf 2 61s
```

Updated Pod definition

```
# Two containers
2  # Container one - talkingcow-env
3  # Container two - Nginx server with https certificate

4  
5  apiVersion: v1
6  kind: Pod
7  metadata:
8   name: nginx-https-pod
9  spec:
10   containers:
11   - image: classpathio/talking-cow-env
12   name: talking-cow-container
13  env:
```

```
- name: INTERVAL
    valueFrom:
      configMapKeyRef:
        name: nginx-conf
        key: sleep-interval
- image: nginx:alpine
  name: web-server
  volumeMounts:
  - name: html
    mountPath: /usr/share/nginx/html
    readOnly: true
  - name: config
    mountPath: /etc/nginx/conf.d
    readOnly: true
  - name: certs
    mountPath: /etc/nginx/certs/
    readOnly: true
  ports:
  - containerPort: 80
  - containerPort: 443
volumes:
- name: html
  emptyDir: {}
- name: config
  configMap:
    name: nginx-conf
    items:
    - key: nginx-configuration.conf
      path: https.conf
- name: certs
  secret:
    secretName: nginx-https
```

```
kubectl create -f pod-nginx-configmap-https.yaml
```

```
    kubectl get po
    NAME READY STATUS RESTARTS AGE
    nginx-https-pod 2/2 Running 0 5s
```

```
kubectl exec -it nginx-https-pod -c web-server sh
// #
```

Verify the Server has picked up the Secret

```
1 kubectl port-forward fortune-https 8443:443 &
2 Forwarding from 127.0.0.1:8443 -> 443
3 Forwarding from [::1]:8443 -> 443
4 $ curl https://localhost:8443 -k
```

```
curl https://localhost:8443 -k -v
   * Rebuilt URL to: https://localhost:8443/
   * Trying 127.0.0.1...
   * TCP_NODELAY set
   * Connected to localhost (127.0.0.1) port 8443 (#0)
6 * ALPN, offering h2
   * ALPN, offering http/1.1
8 * Cipher selection: ALL:!EXPORT:!EXPORT40:!EXPORT56:!aNULL:!LOW:!RC4:@STRE
   * successfully set certificate verify locations:
       CAfile: /etc/pki/tls/certs/ca-bundle.crt
    CApath: none
* TLSv1.2 (OUT), TLS header, Certificate Status (22):
* TLSv1.2 (OUT), TLS handshake, Client hello (1):
* TLSv1.2 (IN), TLS handshake, Server hello (2):
* TLSv1.2 (IN), TLS handshake, Certificate (11):
   * TLSv1.2 (IN), TLS handshake, Server key exchange (12):
   * TLSv1.2 (IN), TLS handshake, Server finished (14):
   * TLSv1.2 (OUT), TLS handshake, Client key exchange (16):
* TLSv1.2 (OUT), TLS change cipher, Change cipher spec (1):
   * TLSv1.2 (OUT), TLS handshake, Finished (20):
21 * TLSv1.2 (IN), TLS change cipher, Change cipher spec (1):
22 * TLSv1.2 (IN), TLS handshake, Finished (20):
* SSL connection using TLSv1.2 / ECDHE-RSA-AES256-GCM-SHA384
* ALPN, server accepted to use http/1.1
25 * Server certificate:
26 * subject: CN=www.classpath.io
27 * start date: Apr 6 07:44:46 2020 GMT
28 * expire date: Apr 4 07:44:46 2030 GMT
   * issuer: CN=www.classpath.io
30 * SSL certificate verify result: self signed certificate (18), continuing
   > GET / HTTP/1.1
32 > Host: localhost:8443
33 > User-Agent: curl/7.61.1
34 > Accept: */*
36 < HTTP/1.1 200 OK
```

```
kubectl logs nginx-https-pod -c web-server
2020/04/06 14:05:52 [error] 6#6: *1 directory index of "/usr/share/nginx/h
127.0.0.1 - - [06/Apr/2020:14:05:52 +0000] "GET / HTTP/1.1" 403 153 "-" "c
2020/04/06 14:06:30 [error] 6#6: *2 directory index of "/usr/share/nginx/h
127.0.0.1 - - [06/Apr/2020:14:06:30 +0000] "GET / HTTP/1.1" 403 153 "-" "c
127.0.0.1 - - [06/Apr/2020:14:07:19 +0000] "GET / HTTP/1.1" 400 255 "-" "c
127.0.0.1 - - [06/Apr/2020:14:08:44 +0000] "GET / HTTP/1.1" 400 255 "-" "c
```

Exposing Secrets through Environment Variables

Instead of configMapKeyRef use secretKeyRef

```
1 env:
2  - name: FOO_SECRET
3    valueFrom:
4    secretKeyRef:
5    name: fortune-https
6    key: foo
```

! Exposing secrets via Env variables is not a best practice as it can be logged in the application logs. Also, the child process will inherit all the environment variables of the parent process making the data vulnerable to leak.

Image Pull Secrets

- When using K8s to pull images from private repositories
- To pull images from private dockerhub registry, there are two steps
 - Create a Secret holding the Dockerhub credentials
 - Reference the secret in the imagePullSecrets field in the mod manifest
- 1. Creaing a image secret

Instead of generic secret use the docker-registry Secret

```
$ kubectl create secret docker-registry mydockerhubsecret \
    --docker-username=myusername --docker-password=mypassword \
    --docker-email=my.email@provider.com
```

kubectl describe mydockerhubsecret

- Contains a single entry called .dockercfg which is equivalent to .dockercfg inside the home directory, which got created when running the docker login command
- 2. Using the Secret in the Pod definition

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4  name: private-pod
5 spec:
6  imagePullSecrets:
7  - name: mydockerhubsecret
8  containers:
9  - image: username/private:tag
10  name: main
```



Deployment

Types of deployment

- Delete all the existing pods and start the new pods
 - hrief downtime

0

- Start all the new pods. Once they are up, delete the old pods
 - Two versions of pods for brief time
 - Consume more resources
 - The new version should not break the schema/contract for the old pods

Deployment using Kubectl

- · Old way of deployment
- Now obsolete and should not be used for newer deployment
- Deployment is driven by the client and not driven through the master
- Inconsistent state of Pods and ReplicationControllers on the server, if there is an issue (network failures)
- Its imperative
- The process should be closely monitor till the kubectl command is successfully executed.

Deployment using Deployment Resource

- Higher level resource
- On top of ReplicaSet
- Creates a ReplicatSet in the background
- When using deployment, actual pods and created and managed by the Replicasets of Deployment resource and not by the Deployment directly
- Coordiantes between old version of ReplicaSet and new version of ReplicaSet
- There will be no change in the Service.

Creating a Deployment Resource

```
1 apiVersion: apps/vlbetal
2 kind: Deployment
3 metadata:
4    name: kubia
5 spec:
6    replicas: 3
7    template:
8    metadata:
9     name: kubia
10    labels:
11    app: kubia
12    spec:
13    containers:
14    - image: luksa/kubia:vl
15    name: nodejs
```

Ensure deleting all the ReplicationController

```
$ kubectl delete rc --all
```

Create the deployment

```
$ kubectl create -f kubia-deployment-v1.yaml --record
```

Note: Make sure you pass in the record argument

Verify the Pre deployment configuration for Pods, RS

Check the deployment status

```
$ kubectl rollout status deployment kubia
deployment kubia successfully rolled out
```

Ensure there are three pods running

\$ kubectl get po

Verify the ReplicaSets

\$ kubectl get replicasets

Updating the Deployment

- Modify the Pod template
- Done declaratively
- K8s will transition the system to the new state.

Deployment Strategies

- RollingUpdate
 - Default strategy
 - Removes pod one by one while adding the new pods
 - Application is available throught the deployment process
 - Use this, when the application supports both versions of the app to run parallelly
- Recreate
 - Deletes all the old pods at one and creates new pods
 - Use this, when your deployment does not support multiple parallel versions and old ones to be stopped completely before the new ones are deployed.
 - Application will be unavailable for short period of time

Starting the deployment process

```
$ kubectl patch deployment kubia -p '{"spec": {"minReadySeconds": 10}}'
```

Note: path command is used to update one or few property of a resource without

Triggerring the Rolling update

• To test the status of deployment, continously run the curl command in a terminal.

```
$ while true; do curl http://130.211.109.222; done
```

Update the deployment to the new version

```
$ kubectl set image deployment kubia nodejs=luksa/kubia:v2
```

Verify the deployment by listing the ReplicaSets

```
$ kubectl get rs
```

Rolling back a deployment

Rolls the deployment to the previous version

```
$ kubectl rollout undo deployment kubia
```

- Rolling the deployment to the specific version
 - Control the revision history with revisionHistoryLimit property (default is 2)

```
$ kubectl rollout undo deployment kubia --to-revision=1
```

Pausing the Rollout

- Typically done for canary release
- One or small fractions of pods will be deployed with the new versions

```
$ $ kubectl set image deployment kubia nodejs=luksa/kubia:v4

$ $ kubectl rollout pause deployment kubia
```

Resuming the Rollout

\$ kubectl rollout resume deployment kubia

Kubernetes Architecture

Kubernetes Architecture

Kubernetes cluster is split into two parts

- K8s control plane
- Worker nodes

Control Plane: - Controls and makes the entire cluster function.

Components of Control plane

- Etcd distributed persistent storage
- API server
- Scheduler
- Controller Manager

Components running on worker nodes

- Kubelet
- K8s service proxy (kube-proxy)
- Container runtime (docker, rkt)

Add on components

- K8s Dashboard
- DNS Server
- Ingress Controller
- Container Network Interface

Checking the status of Control plane

MESSAGE

```
4 controller-manager Healthy ok
5 etcd-0 Healthy {"health": "true"}
```

- · K8s system components do not talk to each other
- K8s components only talk to the API server
- API server only talks to the etcd
- None of the components talk to the etcd
- Components modify the cluster state by talking to the API server

Running multiple instances of individual components

- All the components of the worker nodes need to run on the same node
- Components of the Control plane can be split across multiple servers
- More than one instance of Controle plane can be run to ensure HA
- Multiple instance of etcd and API server can be running in parallel on different server
- Only single instance of Scheduler and ControllerManager can be active at any given time keeping others in the stand-by mode

Use of etcd

- All the K8s resources like pods, RC, RS, Secrets, Services are stored in etcd.
- The manifest files survive the API server restarts and failures.
- Its a key-value data store
- Distributed and hence, more than one etcd server can run increasing the availability and performance
- All the components indirectly read and write entries to etcd using API server
- Only API server talks to etcd directly
- When multiple etcd servers are running in a clustered environment, consistency is achieved using consensus algorithm (RAFT)
- Use odd numbers of etcd to set up cluster

K8s stores all the entries of etcd under /registry directory

Use of API server

- Central component which is used by all the components and clients
- Provides CRUD interface for querying and modifying the cluster state in a consistent manner using REST API
- Stores data in the etcd.
- · Provides validation of objects
- Handles optimistic locking on resources
- Authentication and Authorization can be configured through plugins
- Admission control plugin
 - AlwaysPullImages to force to pull the images when pod is created
 - ServiceAccount to apply default service account, if not specefied
 - ResourceQuota Ensures pods in a namespace uses as much CPU and memory which is allocated
- Store the metadata in the etcd after validating the resource objects
- Clients watch for changes by opening an HTTP connection to the API server
- Clients will receive a stream of modifications to the watched objects
- Every time there is an update, the clients will be notified with the updated object

Example: client watching for any changes using the watch flag

kubectl get pods -watch

Scheduler

- Wait for the newly created pod through the API server's watch
- Assign a node to each new pod that does not have a node set
- Update the pod definition through the API server
- The API server then notifies the Kubelet that the pod is scheduled
- Kubelet on the node sees that the pod is scheduled and thus creates and runs the pod's container
- The Kubelet then creates and runs the pod containers
- Selecting the node can be as simple as selecting a random node and can be as complex as selecting through machine learning algorithms

We can run multiple Schedulers in which case, the mentioned in each pod specification

Default Scheduling Alorithm

- List all the available node that the pod can scheduled to
- Prioritize the acceptable node and choose the best one
- In case of conflicting prioritized nodes, use the round-robin algorithm

Controller Managers

- API Server just stores the resources in the ETCD
- Scheduler assigns the Node
- Other active components are needed to make sure to attain the desired state specified in the resources deployed through API server
- This job is performed by ControllerManagers
- Resources are the descriptors while the ControllerManager is the worker component that does the actual job.
- Controllers run a reconciliation loop by using the watch mechanism to be notified when there is a change and also perform a re-list operation to ensure no events are missed out
- Controllers do not talk to each other directly and only connects to the API server through watch mechanism.

List of Controllers

- ReplicationManager
- ReplicaSet
- DaemonSet
- DeploymentController
- NodeController
- ServiceController
- EndpointController
- NamespaceController
- PersistentVolumeController

Kubelet

- Kubelet and Service Proxy both run on the worker node
- Initially registers with the Node by creating a Node resource on the API server
- Contrinously montiro the API server for pods that have been scheduled to the node and start the containers
- Tells the configured container runtime to pull the image from the registry

Service Proxy

- To ensure that the clients can connect to Services defined through the API server
- When a service is backed by more than on pods, the proxy does the load balancing across those pods

Chain of Events of a Deployment Resource

- Kubectl submits the Deployment resource manifest to API Server through a POST request
- The API Server validates the manifest and stores in etcd
- DeploymentController client is notified about the event
- DeploymentController creates a ReplicaSet manifest through API Server
- ReplicatSetController then picks up the ReplicaSet which considers the replica count and manages the next state
- The Scheduler then chooses the best node and assigns the pods the node
- The Kubelet on the worker node then inspects the pod definition and instructs the container runtime to pull the image and run the containers

Securing K8s API Server

- Pods talk to the API server to retrive or change the state of resources which are deployed in the cluster
- To Authenticate with the API server, service accounts are mounted inside the pod

Authentication

- · Authentication plugins will be configured on the API server
- Same holds good for Authorization plugins which are configured on the API server
- After the API Server receives a request, it goes through a list of authentication plugins
- The plugin extracts the username, userld and groups that the user belongs to and passes this information back to the API server
- Once the information is retrieved, the remaining authentication plugins are discarded and continues to the authorization phase
- K8s does not store the user and group information on the server.

Types of Authentication plugins

- Client certificate
- Authentication token passed in as HTTP header
- Basic HTTP authentication

Users

- Actual humans
- Pods (processes inside the containers)
- Users are managed by external systems
- Pods use a mechanism called Service Accounts
- Service Accounts are managed by Service Accounts resources

Groups

- Both humans and Service Accounts can belong to one or more groups
- Groups are useed to grant permissions to serveral users at once instead of granting permissions at individual level

- · Groups are represented by strings
- There are build in groups
 - o system:unauthenticated
 - o system:authenticated
 - o system:serviceaccounts
 - o sysmts:serviceaccounts:namespace

Service Accounts

- Every Pod is mounted with a secret volume
- The pod authenticates with the API server by sending the contents of the file under `var/run/secrets /kubernetest.io/serviceaccount/token
- API server after successfull authentication allows to perform the operation
- Each pod is associated with a service account which represents the identity of the application
- The token field holds the SA authentication token
- The Authentication plugin authenticates the token, extracts the username, groups and Userld information and passes back to the API Server
- API Server then passes these information to Authorization plugin which determine if the app can perform the operation
- Each pod is associated with only one Service Account in the pod's namespace
- Multiple pods can use the same Service Account
- Each namespace contains a default Service Account

Format of service account username

system:serviceaccount:<namespace>:<service account name>

kubectl get sa



Service Account are tied to Authorization

- · SA can be assigned to the pod in the pod's specification
- If not specified, the pod will use the default service account in the namespace
- different SA can be assigned to different pods

Creating Service Account

- Create additional SA apart from default SA for greater cluster security
- · Pods that should not access cluster level resources should run under a constrained SA

```
1 kubectl create serviceaccount foo
2 serviceaccount "foo" created
```

```
1 kubectl describe sa foo
2 Name: foo
3 Namespace: default
4 Labels: <none>
5
6 Image pull secrets: <none>
7
8 Mountable secrets: foo-token-qzq7j
9
10 Tokens: foo-token-qzq7j
```

The above command creates a custom secret and associates with the SA

```
$ kubectl describe secret foo-token-qzq7j

...
ca.crt: 1066 bytes
namespace: 7 bytes
token: eyJhbGciOiJSUzI1NiIsInR5cCI6IkpXVCJ9...
```

By default, a pod can mount any Secret it wants.

 Pod's ServiceAccount can be configured to only allow to mount Secrets that are listed as mountable Secrets on the Service-Account.

Assigning a SA to Pod

- In the Pod manifest set the SA in the `spec.serviceAccountName field
- SA can be set when creating the pod.
- SA cannot be changed later

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4    name: curl-custom-sa
5 spec:
6    serviceAccountName: foo
7    containers:
8    - name: main
9    image: tutum/curl
10    command: ["sleep", "9999999"]
11    - name: ambassador
12    image: luksa/kubectl-proxy:1.6.2
```

Confirm that the SA secret is mounted in the Pod

```
$ $ kubectl exec -it curl-custom-sa -c main cat /var/run/secrets/kubernetes.
2 eyJhbGci0iJSUzI1NiIsInR5cCI6IkpXVCJ9...
```

Securing the cluster with RBAC (Authorization Plugin)

- System wide Plugin
 - Clients send GET , POST , PUT , DELETE HTTP request
 - Resources are Pods, Service, Secrets etc.
 - Specify the permissions for Standard resources and non-resource URLs
 - RBAC rules are configured throught 4 resources
 - Role and ClusterRoles
 - RoleBinding and ClusterRoleBinding

- Role and RoleBinding are resource level resource
- ClusterRole and ClusterRoleBinding are cluster level resource
- Multiple Roles and RoleBindings can be created in a namespace
- 1. Create namespaces and run the pods

```
$ kubectl create ns foo
namespace "foo" created

kubectl run test --image=luksa/kubectl-proxy -n foo
deployment "test" created

kubectl create ns bar
namespace "bar" created

kubectl run test --image=luksa/kubectl-proxy -n bar
deployment "test" created
```

2. Fetch the name of the pod

```
1 $ kubectl get po -n foo
2 NAME READY STATUS RESTARTS AGE
3 test-145485760-ttq36 1/1 Running 0 1m
```

3. Login to the pod

```
1 $ kubectl exec -it test-145485760-ttq36 -n foo sh
2 / #
```

4. Perform the above steps for the other pod as well

Create a Role resource

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
namespace: foo
name: service-reader
rules:
- apiGroups: [""]
verbs: ["get", "list"]
resources: ["services"]
```

! plural form should be used when specifying the resources

```
1 $ kubectl create -f service-reader.yaml -n foo
2 role "service-reader" created
```

Using the command directly to create the role

```
1 $ kubectl create -f service-reader.yaml -n foo
2 role "service-reader" created
```

Binding the Role to Service Account

- Role defines what actions can be performed
- RoleBinding specefies who can perform
- RoleBinding can bind a role to user, SA or a group
- RoleBinding always referes to a single role
- A Role can bind to multiple subjects (users, groups, sevice accounts)

•

Creating a RoleBinding resource

```
$ kubectl create rolebinding test --role=service-reader --serviceaccount=f
2 rolebinding "test" created
```

Describe the RoleBinding

```
kubectl get rolebinding test -n foo -o yaml
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
    name: test
    namespace: foo
    ...
roleRef:
    apiGroup: rbac.authorization.k8s.io
    kind: Role
    name: service-reader
subjects:
    - kind: ServiceAccount
name: default
namespace: foo
```

Include the SA from another namespace in Rolebinding

Edit the existing Rolebinding

```
$ kubectl edit rolebinding test -n foo
```

Reference another SA in the current RoleBinding

```
1 subjects:
2 - kind: ServiceAccount
3    name: default
4    namespace: bar
```

ClusterRole and ClusterRoleBinding

- Role and RoleBinding are namespaced resource
- Clusterlevel resources are managed by ClusterRole and ClusterRoleBinding
- To manage non-resource URL
- The default SA cannot list the PV, even after the clusterRole has bound to SA

```
1 kubectl create clusterrole pv-reader --verb=get,list --resource=persisten2 clusterrole "pv-reader" created
```

```
kubectl get clusterrole pv-reader -o yaml
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
name: pv-reader
resourceVersion: "39932"
selfLink: ...
uid: e9ac1099-30e2-11e7-955c-080027e6b159
rules:
- apiGroups:
- ""
resources:
- persistentvolumes
verbs:
- get
- list
```

```
kubectl create rolebinding pv-test --clusterrole=pv-reader --serviceaccourolebinding "pv-test" created
```

To access cluster level resources from the service account, bind the ClusterRoleBinding

Delete the RoleBinding

```
1 $ kubectl delete rolebinding pv-test
```

2 rolebinding "pv-test" deleted

Create ClusterRoleBinding

```
    $ kubectl create clusterrolebinding pv-test --clusterrole=pv-reader
    --serviceaccount=foo:default
    clusterrolebinding "pv-test" created
```

Allowing to access non-resource URL's

Predefined ClusterRoles and ClusterRoleBinding

Cluster role system:discovery

```
$ $ kubectl get clusterrole system:discovery -o yaml
  apiVersion: rbac.authorization.k8s.io/v1
3 kind: ClusterRole
4 metadata:
   name: system:discovery
   . . .
7 rules:
8 - nonResourceURLs:
    - /api
   - /api/*
   - /apis
   - /apis/*
   - /healthz
   - /swaggerapi
   - /swaggerapi/*
   - /version
   verbs:
   - get
```

ClusterRoleBinding for system:discover role

```
1 $ kubectl get clusterrolebinding system:discovery -o yaml
2 apiVersion: rbac.authorization.k8s.io/v1
```

```
kind: ClusterRoleBinding
metadata:
name: system:discovery
...
roleRef:
apiGroup: rbac.authorization.k8s.io
kind: ClusterRole
name: system:discovery
subjects:
- apiGroup: rbac.authorization.k8s.io
kind: Group
name: system:authenticated
- apiGroup: rbac.authorization.k8s.io
kind: Group
name: system:authorization.k8s.io
kind: Group
name: system:unauthenticated
```

Cluster Role view

```
$ $ kubectl get clusterrole view -o yaml
  apiVersion: rbac.authorization.k8s.io/v1
3 kind: ClusterRole
  metadata:
   name: view
7 rules:
8 - apiGroups:
   resources:
   configmaps
   endpoints

    persistentvolumeclaims

   - pods
   - replicationcontrollers
   - replicationcontrollers/scale
   - serviceaccounts
   - services
   verbs:
   - get
   - list
   - watch
```

• Cluster Role can be bound to ClusterRoleBinding or RoleBinding

- ClusterRoleBinding bound to ClusterRole, the subjects listed in the binding can access resources across all namespaces
- If RoleBinding is bound to ClusterRole, the subjects listed in the binding can access only namespaced resources

Predefined ClusterRoles and ClusterRoleBindings

kubectl get clusterrolebindings	
NAME	AGE
cluster-admin	1d
system:basic-user	1d
system:controller:attachdetach-controller	1d
•••	
system:controller:ttl-controller	1d
system:discovery	1d
system:kube-controller-manager	1d
system:kube-dns	1d
system:kube-scheduler	1d
system:node	1d
system:node-proxier	1d
kubectl get clusterroles	
NAME	AGE
admin	1d
cluster-admin	1d
edit	1d
system:auth-delegator	1d
system:basic-user	1d
system:controller:attachdetach-controller	1d
•••	
system:controller:ttl-controller	1d
system:discovery	1d
system:heapster	1d
system:kube-aggregator	1d
system:kube-controller-manager	1d
system:kube-dns	1d
system:kube-scheduler	1d
system:node	1d
system:node-bootstrapper	1d
system:node-problem-detector	1d
system:node-proxier	1d
system:persistent-volume-provisioner	1d
view	1d

• Important roles are view edit admin and cluster-admin

View

- · Reading resources in a namespace
- Except Roles, RoleBindings and secrets

Edit

- Modify resources including reading and modifying secrets
- Forbids reading and modifying Roles and RoleBindings

Admin

- Complete control of resources in the namespace including Roles and RoleBindings
- Except ResourceQuota

Cluster-Admin

• Complete contol of resources including Resource Quota

Advance Scheduling

Taints and Tolerations to repel pods

- Node taints and Pods tolerations
- Used to restrict which pods can use a certain node
- Pod can only be scheduled to a node if it can tolerate the Node taints
- Nodes can have more than 1 taint and Pods can have more than 1 tolreations
- Can be used for Cluster partitioning to dev.preprod etc

```
kubectl describe node master.k8s

Name: master.k8s

Role:
Labels: beta.kubernetes.io/arch=amd64
beta.kubernetes.io/os=linux
kubernetes.io/hostname=master.k8s
node-role.kubernetes.io/master=
Annotations: node.alpha.kubernetes.io/ttl=0
volumes.kubernetes.io/controller-managed-attach-detach=true
node-role.kubernetes.io/master:NoSchedule
...
```

Node Taints

- · Contains a Key, Value and Effect
- Represented by <key>=<value>:<effect>
- Three kinds of effects
 - NoSchedule Pods would not be scheduled if they cannot tolerate the taint
 - PreferNoSchedule Softer version of NoSchedule Avoind scheduling pods if they cannot tolerat the taint
 - NoExecute The running pods are also evicted if they cannot tolerate the taints

Pod Tolearation

```
    kubectl describe po kube-proxy-80wqm -n kube-system
    ...
    Tolerations: node-role.kubernetes.io/master=:NoSchedule
```

```
node.alpha.kubernetes.io/notReady=:Exists:NoExecute
node.alpha.kubernetes.io/unreachable=:Exists:NoExecute
...
```

Custom Taints

Use case - Creating a Taint on production Node to avoid pre-prod pods to be scheduled on Prod Node

```
1 kubectl taint node node1.k8s node-type=production:NoSchedule
2 node "node1.k8s" tainted
```

```
kubectl run test --image busybox --replicas 5 -- sleep 99999
   deployment "test" created
   kubectl get po -o wide
5 NAME
                     READY STATUS
                                                              NODE
                                    RESTARTS
                                              AGE IP
6 test-196686-46ngl 1/1
                                              12s
                                                   10.47.0.1
                                                              node2.k8
                           Running
                                              12s 10.47.0.7
                                                              node2.k8
7 test-196686-73p89 1/1
                           Running 0
                                              12s 10.47.0.6
8 test-196686-77280 1/1
                           Running 0
                                                              node2.k8
9 test-196686-h9m8f 1/1
                           Running 0
                                              12s 10.47.0.5
                                                              node2.k8
                                                              node2.k8
10 test-196686-p85ll 1/1
                           Running
                                              12s
                                                   10.47.0.4
```

Adding Toleartions to Pods

Pod Manifest

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
name: prod
spec:
replicas: 5
template:
spec:

tolerations:
key: node-type
```

```
operator: Equal
value: production
effect: NoSchedule
```

```
$ kubectl get po -o wide
  NAME
                    READY STATUS
                                   RESTARTS
                                             AGE IP
                                                             NODE
                                                             node1.k8
3 prod-350605-1ph5h 0/1
                          Running
                                             16s 10.44.0.3
4 prod-350605-ctqcr 1/1
                                             16s 10.47.0.4
                          Running
                                                             node2.k8
                                   0
5 prod-350605-f7pcc 0/1
                          Running 0
                                             17s 10.44.0.6
                                                             node1.k8
6 prod-350605-k7c8g 1/1
                                             17s 10.47.0.9
                                                             node2.k8
                          Running
                                   0
  prod-350605-rp1nv 0/1
                          Running
                                   0
                                             17s 10.44.0.4
                                                             node1.k8
```

Node Affinity to attract Pods to certain Nodes

• Allows to tell K8s to schedule pods only to specific subset of nodes

Node Affinity vs Node Selectors

Node Selector	Node Affinity
Old way of affinity mechanism	Newer way of specifying Node Affinity
Node has to include all the labels specified to be eligible to become target for the pod	Can give hard requirements or preferences
Will eventually be deprecated	New standard going further

Important Node labels

- failure-domain.beta.kubernetes.io/region specifies the geographical region
- failure-domain.beta.kubernetes.io/zone specifies the zone
- kubernetes.io/hostname specifies the hostname

Specifying the Node Affinity in the Pod manifest

More expressive than NodeSelector

```
apiVersion: v1
kind: Pod
metadata:
name: kubia-gpu
spec:
affinity:
nodeAffinity:
requiredDuringSchedulingIgnoredDuringExecution:
nodeSelectorTerms:
- matchExpressions:
- key: gpu
operator: In
values:
- "true"
```

- requiredDuringScheduling Rules defined under this specify the labels the node must have for the pod to be scheduled
- ...IgnoredDuringScheduling Rules defined under this will be ignored during scheduling and dont affect the pods already executing on the node

Prioritizing the nodes while scheduling a Pod

 Node affinity allows to specify which node the scheduler should prefer when scheduling a pod using the preferredDuringSchedulingIgnoredDuringExecution field

Label the Nodes appropriately

• Based on Availability zone

```
kubectl label node node1.k8s availability-zone=zone1
node "node1.k8s" labeled

kubectl label node node1.k8s share-type=dedicated
node "node1.k8s" labeled

kubectl label node node2.k8s availability-zone=zone2
node "node2.k8s" labeled

kubectl label node node2.k8s share-type=shared
```

```
node "node2.k8s" labeled
kubectl get node -L availability-zone -L share-type
NAME
             STATUS
                       AGE
                                 VERSION
                                            AVAILABILITY-ZONE
                                                                SHARE-TYPE
master.k8s
             Ready
                       4d
                                  v1.6.4
                                            <none>
                                                                 <none>
node1.k8s
             Ready
                       4d
                                 v1.6.4
                                            zone1
                                                                dedicated
node2.k8s
             Ready
                       4d
                                 v1.6.4
                                            zone2
                                                                 shared
```

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: pref
spec:
  template:
    spec:
      affinity:
        nodeAffinity:
          preferredDuringSchedulingIgnoredDuringExecution:
          - weight: 80
            preference:
              matchExpressions:
              - key: availability-zone
                operator: In
                values:
                 - zone1
          - weight: 20
            preference:
              matchExpressions:
              - key: share-type
                operator: In
                values:
                 - dedicated
```

```
kubectl get po -o wide
NAME
                    READY
                            STATUS
                                      RESTARTS
                                                AGE
                                                      ΙP
                                                                  NODE
pref-607515-1rnwv
                   1/1
                            Running
                                      0
                                                4m
                                                      10.47.0.1
                                                                  node2.k8
pref-607515-27wp0
                  1/1
                            Running
                                     0
                                                4m
                                                      10.44.0.8
                                                                  node1.k8
pref-607515-5xd0z
                  1/1
                            Running
                                     0
                                                4m
                                                      10.44.0.5
                                                                  node1.k8
pref-607515-jx9wt
                  1/1
                                                      10.44.0.4
                                                                  node1.k8
                            Running
                                      0
                                                4m
pref-607515-mlgqm
                  1/1
                            Running
                                      0
                                                4m
                                                      10.44.0.6
                                                                  node1.k8
```

Colocating Pods with affinity and Anti affinity

· To definit the affinity between pods

Deploy one pod

```
1 $ kubectl run backend -l app=backend --image busybox -- sleep 999999
2 deployment "backend" created
```

Front end pod definition

```
1 apiVersion: extensions/vlbeta1
2 kind: Deployment
3 metadata:
4 name: frontend
5 spec:
6 replicas: 5
7 template:
8 ...
9 spec:
10 affinity:
11 podAffinity:
12 requiredDuringSchedulingIgnoredDuringExecution:
13 - topologyKey: kubernetes.io/hostname
14 labelSelector:
15 matchLabels:
16 app: backend
17 ...
```

Deploying pods with Affinity

```
1 $ kubectl get po -o wide
2 NAME READY STATUS RESTARTS AGE IP NODE
3 backend-257820-qhqj6 1/1 Running 0 8m 10.47.0.1 node2.k8s
```

```
1 kubectl create -f frontend-podaffinity-host.yaml
```

```
deployment "frontend" created
   kubectl get po -o wide
   NAME
                        READY
                              STATUS
                                        RESTARTS AGE IP
                                                                NODE
                                                                node2.k8
6 backend-257820-qhqj6 1/1
                               Running
                                                 8m
                                                      10.47.0.1
   frontend-121895-2c1ts 1/1
                               Running
                                        0
                                                 13s 10.47.0.6
                                                                node2.k8
8 frontend-121895-776m7 1/1
                                                 13s 10.47.0.4
                                                                node2.k8
                               Running
                                        0
   frontend-121895-7ffsm 1/1
                               Running
                                        0
                                                 13s 10.47.0.8 node2.k8
10 frontend-121895-fpgm6 1/1
                                                 13s 10.47.0.7 node2.k8
                               Running
                                       0
11 frontend-121895-vb9ll 1/1
                               Running
                                        0
                                                 13s 10.47.0.5
                                                                node2.k8
```

Co-locating pods in the same Availability Zone / Region

```
    Set the toplogyKey to failure-domain.beta.kubernetes.io/zone
    Set the toplogyKey to failure-domain.beta.kubernetes.io/region
```

Expressing Affinity rather than hard requirements

- · Expressing preferences instead of hard requirements
- The scheduler will honor if there is provision else will select a different node

```
apiVersion: extensions/v1beta1
  kind: Deployment
3 metadata:
   name: frontend
  spec:
    replicas: 5
    template:
      . . .
      spec:
         affinity:
           podAffinity:
             preferredDuringSchedulingIgnoredDuringExecution:
             - weight: 80
               podAffinityTerm:
                 topologyKey: kubernetes.io/hostname
                 labelSelector:
                   matchLabels:
                     app: backend
         containers: ...
```

Scheduling the pods with Anti-affinity

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
 name: frontend
spec:
  replicas: 5
  template:
    metadata:
      labels:
        app: frontend
    spec:
      affinity:
        podAntiAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
          - topologyKey: kubernetes.io/hostname
            labelSelector:
              matchLabels:
                app: frontend
      containers: ...
```

```
$ kubectl get po -l app=frontend -o wide
  NAME
                         READY STATUS RESTARTS AGE IP
                                                                NODE
3 frontend-286632-0lffz
                         0/1
                               Pending 0
                                                 1m
                                                      <none>
  frontend-286632-2rkcz
                         1/1
                               Running 0
                                                      10.47.0.1
                                                                node2.k8
                                                 1m
5 frontend-286632-4nwhp
                         0/1
                               Pending 0
                                                 1m
                                                      <none>
6 frontend-286632-h4686
                         0/1
                               Pending 0
                                                 1m
                                                      <none>
  frontend-286632-st222
                         1/1
                               Running 0
                                                      10.44.0.4
                                                                node1.k8
                                                 1m
```

Appendix-A

Ingress Controller Setup

1. Instructions for setup - https://docs.nginx.com/nginx-ingress-controller/installation/installation-with-manifests/

2.

a117bf1ffc495407cb9dbc76bed4ef73-79 2800871.ap-south-1.elb.amazonaws.com 13.234.250.206 13.126.122.135

curl -resolve cafe.example.com:\$IC_HTTPS_PORT:\$IC_IP https://cafe.example.com:\$IC_HTTPS_PORT/tea -insecure

curl –resolve hostnameservice.classpath.com:443:13.126.122.135 https://hostnameservice.classpath.com:443/ –insecure

curl –resolve cafe.example.com:443:13.126.122.135 https://cafe.example.com:443/coffee –insecure

curl –resolve hostnameservice.classpath.com:80:13.126.101.254 http://hostnameservice.classpath.com:80/ –insecure